FOLLOW MY VOICE: SELF-SIMILAR AND PREFERRED VOICES FOR GAZE-CONTINGENT OBJECT SEARCH IN MIXED REALITY

by

PEDRO POVEDA QUEVEDO

A thesis submitted in partial fulfilment of the requirements
for the degree of Master of Science
in the Department of Computer Science
in the College of Engineering and Computer Science
at the University of Central Florida
Orlando, Florida

Fall Term
2026

© 2026 Pedro Poveda Quevedo

GENERATIVE ARTIFICIAL INTELLIGENCE USE DISCLOSURE

This first-pass manuscript was prepared with assistance from OpenAI Codex for source organization, drafting, implementation-to-protocol reconciliation, and LaTeX production. Claude performed separate language and methods reviews; their findings and the changes made in response are recorded with the source project. The author must review the manuscript, verify its citations and claims, and approve the final disclosure before submission. This working disclosure does not assert that author validation or committee approval has already occurred.

ABSTRACT

Spoken guidance can help mixed-reality users locate objects outside their current field of view. Voice identity may also affect how users experience that assistance. Follow My Voice compares a participant-preferred catalog voice with a synthetic version of the participant's own voice. Target identity is withheld while gaze-contingent directions guide search across eight surrounding vertical planes. A right-controller ray and trigger provide explicit selection.

The system separates local spatial reasoning from constrained language-model decisions. A central guidance state machine combines valid gaze, controller intersections, and recent history. Jev selects eligible phrasing and speaking opportunities; local playback checks reject stale cues. Prepared ElevenLabs audio supplies both conditions. Synchronized records distinguish gaze contact, spoken acknowledgment, controller pointing, and successful selection.

The proposed within-participant study comprises one practice trial in the first assigned voice and two counterbalanced blocks of 40 measured trials. Analysis will compare active completion time while accounting for starting geometry and participant clustering. Workload, motivation, and agent experience will be assessed retrospectively after both blocks. This first pass documents the system, measurement definitions, planned analyses, and remaining validation requirements. Confirmatory participant outcomes are not yet available. The contribution at this stage is an auditable experimental platform and a design for studying personalized voice within a shared guidance mechanism.

PREFACE

This is an advisor-review draft dated October 22, 2026. It expands the project's current Overleaf manuscript and bibliography into a thesis using the supplied UCF ETD LuaLaTeX template. It is not a completed empirical report. Present-tense system descriptions refer to the documented implementation snapshot; future-tense methods describe the proposed study. Engineering checks, planned measurements, and participant outcomes are kept separate.

The source snapshot is Overleaf revision `f9f49e4`. The runtime alignment audit dated October 20 is the accompanying implementation record. Later software changes require a new reconciliation before submission. Exact pre-cue geometry, balanced vertical allocation, final sample size, and pilot acceptance rules remain explicit readiness requirements. Qualtrics synchronization remains on hold. The author name, degree details, graduation term, previous degree, advisor, acknowledgments, and final AI disclosure require author confirmation; unverified prior degrees and advisor names have not been inserted.

The structural reference is Hoang Bao Truong's 2026 UCF thesis, *Using Inferences from Natural Language Narratives to Improve Video Question Answering*. Its chapter hierarchy, front matter, detailed methods, diagnostic discussion, and appendices inform presentation only. Its prose, experiments, and findings are not evidence for this study.

TABLE OF CONTENTS

LIST OF FIGURES xiii

LIST OF TABLES xiv

CHAPTER 1: INTRODUCTION 1

Motivation and Background 1

Problem Statement 2

Research Objectives and Questions 3

Scope of the Comparison 4

Contributions of This First Pass 5

Organization of the Thesis 6

CHAPTER 2: LITERATURE REVIEW 7

Visual Search and Attention Guidance in Extended Reality 7

Gaze-Contingent Auditory Guidance 8

Voice as a Social and Personalized Interface 9

Self-Similar Voice and Embodied Collaboration 10

Naturalness, Familiarity, and Uncanniness 11

Context and Short-Term Memory in XR Agents 11

Eye-Tracking Quality and Interpretation 12

Workload, Motivation, and Agent Perception 13

Synthesis and Research Gap 14

CHAPTER 3: SYSTEM DESIGN AND IMPLEMENTATION 15

Apparatus 15

Search Task 15

Sensing and Guidance State 17

Short-Term Interaction Memory 17

Agent Decisions and Playback Control 18

Voice Preparation and Telemetry 20

Trial Review Interface 21

Ownership and State Transitions 21

A Worked Guidance Sequence 23

Asynchronous Decisions and Cancellation 25

CHAPTER 4: EXPERIMENTAL METHODOLOGY 26

Task 26

Design and trial allocation 27

Measures 29

Task performance 29

Starting geometry and search behavior 30

Subjective measures 32

Measurement quality 35

Research questions 35

Participants 35

Procedure 36

Pilot testing and parameter checks 38

CHAPTER 5: MEASUREMENT AND DATA PROVENANCE 39

From Interaction Events to Research Variables 39

Coordinate Frames and Starting Geometry 39

Trial Clocks and Pause Intervals 42

Gaze Contact, Vicinity, Pointing, and Selection 44

Acquisition Rate, Validity, and Gaps 45

Motion, Coverage, and Instructions 46

Sequence-Level Latency 47

Questionnaire Joins and Versioning 49

CHAPTER 6: PLANNED STATISTICAL ANALYSIS 51

Primary contrast 51

Secondary and exploratory outcomes 51

Incomplete and interrupted observations 53

Guidance and qualitative evaluation 54

A Candidate Primary Model 54

Sample Size and Precision Planning 55

Data Screening and Missingness 56

Secondary Outcomes and Multiplicity 57

Behavior as Outcome Rather Than Automatic Adjustment 58

Qualitative Analysis and Mixed Evidence 58

Safety Outcomes and Manipulation Checks 59

CHAPTER 7: ENGINEERING VERIFICATION AND RESULTS FRAMEWORK 60

Evidence Available for This Draft 60

Why Sequence Evaluation Is Necessary 60

Paired Policy Pilots 61

Metrics and Episode Definitions 62

Illustrative Failure Analysis 63

Participant Results to Be Added 64

Interpreting Two-Second Delays 65

CHAPTER 8: DISCUSSION AND LIMITATIONS 66

What the Study Can Explain 66

Possible Outcome Patterns 66

Voice Identity and Acoustic Confounds 67

Adaptive Guidance as a Shared Mechanism 68

Practice, Learning, and Order 69

Retrospective Questionnaires and Burden 69

Gaze Quality and Feedback Coupling 70

Privacy, Consent, and Reproducibility 71

Future Work 71

CHAPTER 9: CONCLUSION 73

APPENDIX A: QUESTIONNAIRE WORDING AND ADMINISTRATION 74

Administration 75

Baseline 76

Current symptoms 77

Voice setup 78

Final survey: condition-specific questionnaire sets 80

Workload 80

Activity experience: IMI 81

Agent perception 82

Voice and assistance experience 84

Technical and comfort checks 85

Final survey: symptoms, comparison and debrief 86

Post-exposure symptoms and comfort 86

Voice comparison and debrief 87

Retained Voice-Recording Passage 89

APPENDIX B: DATA CONTRACT AND SCORING RULES 90

Record Identity and Provenance 91

Geometry and Timing Fields 92

Raw Observation Requirements 93

Scoring Contract 94

Export Validation Cases 94

APPENDIX C: GUIDANCE SEQUENCE TESTS AND EVALUATION RUBRIC 96

Controlled Scenario Matrix 97

Deterministic Labels 98

Advisory Model Review 99

Reporting Template 99

APPENDIX D: STUDY READINESS AND SOURCE REGISTER 101

Implemented, Proposed, and Unresolved 102

Source Register 103

Figure Provenance 104

Submission Work Remaining 104

LIST OF REFERENCES 105

LIST OF FIGURES

Figure 1.1: Voice comparison within the shared search task 5

Figure 3.1: Search-scene illustration 16

Figure 3.2: Bounded spatial and speech memory 19

Figure 3.3: Guidance architecture and playback control 20

Figure 3.4: Trial Review interface with synthetic demonstration 22

Figure 3.5: Guidance states and evidence-dependent transitions 24

Figure 4.1: Study procedure and trial timing 28

Figure 4.2: Final questionnaire timing and block reference scope 37

Figure 5.1: Raw records, derived summaries, and analytical scope 40

Figure 5.2: Nominal wall assignment and exact starting angle 42

Figure 5.3: Cross-relative height and live gaze refinement 43

Figure 5.4: Gaze acknowledgment, pointing acknowledgment, and completion 45

Figure 5.5: Qualification-to-playback latency and its boundaries 48

Figure 6.1: Trial nesting and questionnaire analysis units 52

LIST OF TABLES

Table 3.1: Guidance responsibilities and evidence boundaries. 23

Table 4.1: Study allocation 28

Table 4.2: Condition-specific questionnaire battery 34

Table 5.1: Timing variables and evidence boundaries 49

Table 6.1: Planned outcomes and analysis units 53

Table 7.1: Evidence status at the documented implementation checkpoint. 61

Table 7.2: Planned empirical results, not observed findings. 64

Table B.1: Required record identity and scope. 91

Table B.2: Geometry and event contract. 92

Table C.1: Sequence scenarios and acceptance criteria. 97

Table D.1: Decisions required before confirmatory collection. 102

CHAPTER 1: INTRODUCTION

Motivation and Background

A spoken direction in mixed reality has both spatial content and a speaker. “Look to your right” describes an action, but the voice that delivers it can also sound familiar, reassuring, distracting, or unfamiliar. An assistance system must make the instruction useful at the moment it is heard. It must also account for the fact that the same words may be experienced differently when spoken by a preferred synthetic voice or by a synthetic voice resembling the listener.

Follow My Voice studies this relationship in a surrounding object-search task. A seated participant wears a mixed-reality headset and views colored geometric objects distributed across eight vertical planes. The application knows which object is the target, while the participant is not given its color or shape at the start. The participant instead follows spoken directions, looks for the target, aims a right-controller ray, and presses the trigger to select it. This division makes assistance central to the task without making gaze itself a selection command.

Prior work has investigated gaze-contingent auditory feedback for visual search, multistage gaze guidance, and cues for objects beyond the field of view [25, 21, 5]. Separately, self-similar voices have been studied in educational games and collaboration with virtual characters [19, 16]. These lines of work motivate a question about their intersection: when the assistance mechanism is shared, does the identity of the voice change the efficiency or experience of being guided toward an initially unnamed target?

The present comparison uses two personalized voices. In one condition, the participant selects a preferred voice from a catalog. In the other, the voice is synthesized from a recording of the participant. A preferred catalog voice is an active comparison, not an intentionally poor or impersonal

baseline. The design therefore asks whether self-similar synthesis offers something different from an alternative the participant has already chosen to hear.

This distinction matters for interpretation. If a self-similar voice receives higher similarity ratings, that establishes a manipulation check, not a performance benefit. If participants prefer it but do not search faster, the preference remains an experience result. If they search faster without preferring it, behavior and liking have diverged. The thesis will treat these possibilities as separate outcomes rather than combine them into a single claim that one agent is better.

Problem Statement

The research problem has an empirical component and an engineering component. The empirical component is to compare voice conditions without changing the guidance policy, target-selection mechanism, or intended difficulty distribution. The engineering component is to ensure that adaptive speech remains appropriate while the participant's gaze and controller move. A direction that was correct when requested can become obsolete before playback. A controller may already point at the target while an older downward correction is still waiting to speak. These failures can influence both search time and perceived responsiveness, thereby contaminating a voice comparison.

Gaze measurements also require careful interpretation. A ray intersecting an object is an observation made by the interaction system. It does not establish that the participant fixated, recognized, or understood the object. A nearby-gaze region is a deliberate tolerance used by the interface. It may improve usability, but it is not a substitute for measuring spatial accuracy. Similarly, a headset's nominal output rate does not demonstrate that the application receives distinct gaze observations at that rate [12, 30].

The thesis addresses these problems by separating observed evidence, guidance state, model decisions, and delivered speech. The same distinctions carry into the analysis. An instruction selected by a model is not counted as spoken until playback starts. A qualified gaze acknowledgment is not a controller acknowledgment. A controller acknowledgment is not a successful trial until the trigger selection is correct. This separation is necessary both to implement responsive guidance and to explain what the participant actually experienced.

Research Objectives and Questions

The primary objective is to evaluate a voice-condition contrast within a shared gaze-contingent assistance system. The primary research question is: will active search-completion time differ between self-similar and preferred catalog voices after accounting for absolute starting angle and repeated observations within participants? The contrast is two-sided because the current design does not justify assuming a performance advantage for self-similarity.

Secondary questions concern perceived workload and activity experience. Will the two voices differ on the six workload dimensions and their complete unweighted mean? Will they differ on interest/enjoyment, effort/importance, pressure/tension, or value/usefulness? These outcomes concern different constructs. The analysis will preserve those distinctions rather than pool them into a broad satisfaction score.

Exploratory questions concern agent perception, reliance, and the relationship between task difficulty and voice experience. Participants may describe a voice as more responsive or humanlike, may rely on guidance more often, or may experience a cloned voice as familiar or eerie. Such reports will be treated as judgments about an interaction. They will not be evidence that the system possesses intelligence, emotions, awareness, or an inner voice.

An additional engineering objective is to make the assistance process inspectable. Can recorded state and speech history explain why a cue was eligible, why it played or was suppressed, and how long it waited? Can evaluation distinguish a genuinely unhelpful instruction from an acceptable alternative wording? These questions concern system validity and readiness. They support the empirical study but do not replace it.

Scope of the Comparison

The independent variable is voice condition: participant-preferred catalog voice versus self-similar synthetic voice. Both conditions use gaze contingency, the same semantic phrase library, the same local geometry rules, and the same controller-selection mechanics. Consequently, the study will not isolate the benefit of gaze contingency, short-term memory, the cloud decision agent, or synthetic speech in general. Those are shared system properties.

The task is deliberately narrower than unrestricted navigation. Participants remain seated and rotate their head and chair to search surrounding planes. Guidance such as "turn left" concerns orientation within this setup; it does not ask them to walk toward a target. Objects are virtual, the physical room is visible through passthrough, and the target is initially unnamed. Conclusions will therefore be bounded to guided search under these conditions.

The use of one practice trial in the first assigned voice is a current protocol decision. It introduces an exposure asymmetry: whichever voice appears first also appears during practice. Counterbalancing distributes that asymmetry across voices, but does not remove learning or carryover. The analysis must retain period and assigned order, and the report must describe the practice assignment explicitly.

Figure 1.1 separates the voice manipulation from the shared task. A difference between conditions would require evidence from the planned study; the diagram specifies the comparison rather than predicting its direction.

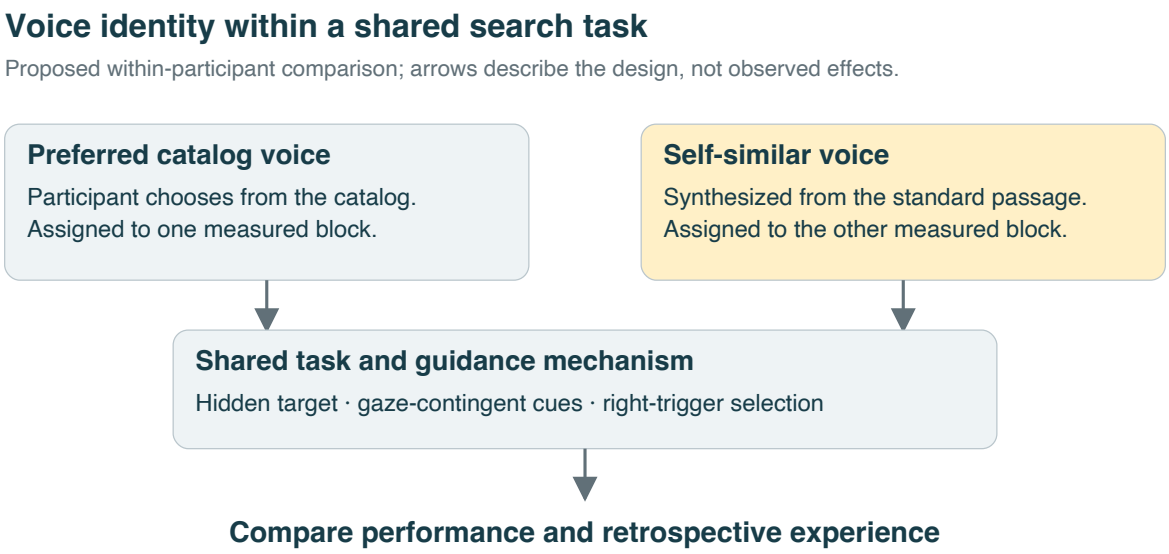


Figure 1.1: Proposed comparison of preferred and self-similar voices within a shared task and guidance mechanism. The schematic identifies the intended manipulation and outcome families; it does not depict observed effects.

Contributions of This First Pass

The implemented contribution is an instrumented mixed-reality guidance platform. It combines local spatial eligibility, a central guidance state machine, bounded history, constrained cloud decisions, prepared speech, and synchronized review. The resulting architecture exposes the relationship between what the system observed and what it delivered.

The methodological contribution is an operational study design. It specifies two counterbalanced voice blocks, a hidden-target task, independent right-trigger selection, trial geometry, event bound-

aries, retrospective questionnaires, and a planned repeated-measures analysis. The measurement definitions are intended to prevent several common substitutions: nominal wall angle for exact starting angle, hover for fixation, proposals for playback, or duplicated block ratings for independent trial observations.

The empirical contribution remains pending. This draft does not report a recruited sample, condition means, significance tests, or participant quotations. Its engineering chapter records what the existing alignment audit supports and defines how later evidence will be reported. A completed thesis will need to replace the pending results sections with the frozen study's actual outcomes and revise the abstract and conclusion accordingly.

Organization of the Thesis

Chapter 2 relates the design to visual-search guidance, voice identity, context-aware agents, and measurement. Chapter 3 describes the headset application and the division of responsibility between Unity and Jev. Chapter 4 specifies the proposed study. Chapter 5 defines geometry, timing, and data provenance, and Chapter 6 explains how those observations will support the research questions. Chapter 7 separates engineering verification from future participant results. The discussion and conclusion identify the claims the design can support and the work still required. Appendices preserve the current instrument wording, data contract, sequence tests, and study-readiness decisions.

CHAPTER 2: LITERATURE REVIEW

Visual Search and Attention Guidance in Extended Reality

Search in a surrounding display differs from search on a single front-facing screen because relevant objects may lie outside the visible field. Turning changes what can be inspected, while eye movements refine attention within the visible region. Work on eye and head movements across an extended field of view is therefore relevant to the task geometry [36]. In the present study, this motivates recording both head motion and gaze rather than treating one as a complete substitute for the other.

The project's bibliography includes several approaches to directing attention in immersive environments. Grogorick et al. address subtle gaze guidance, while HiveFive concerns immersion-preserving attention guidance [14, 22]. Ariza and colleagues investigate vibrotactile assistance and proximity-based multimodal feedback for selection [3, 2]. These works establish relevant design alternatives; they do not imply that all forms of guidance are equivalent. A cue that changes visual salience, a vibration that indicates proximity, and a spoken spatial instruction impose different perceptual and interpretive demands.

Binetti et al. examine visual and auditory cues for locating out-of-view objects in head-mounted augmented reality [5]. Warden et al. focus on target cue type and location in augmented-reality search [37]. Together, these topics make clear why target visibility and cue availability belong in the method description. Follow My Voice does not present an opening color-and-shape target description. Instead, the direction sequence gradually narrows the participant's search.

This hidden-target design creates an important boundary. A participant who already knows the target can continue searching without further assistance. Here, the participant must obtain enough

spatial information to infer which object is intended. The voice may consequently be heard more often and may be more consequential than in a task where it supplements an already known target. That is a reason to describe the task precisely, not a reason to assume stronger voice effects in advance.

Gaze-Contingent Auditory Guidance

Losing et al. connect gaze-contingent auditory feedback with visual-search guidance [25]. Two-Step Gaze Guidance provides another relevant precedent for staged guidance [21]. The present system adopts a coarse-to-fine organization: orientation toward the correct wall precedes within-wall horizontal and vertical refinement. This is a design adaptation. Its particular geometry, phrases, timers, and confirmation regions require validation in the present apparatus.

Zhang et al. examine visual and audio hints alongside gaze-assisted post-task feedback in augmented-reality visual search [40]. That combination is useful for distinguishing guidance during a task from feedback after it. Follow My Voice similarly separates navigation phrases, target acknowledgments, and success feedback in its logs. A success chime after selection cannot explain an earlier change in search direction in the same trial, although it may influence subsequent experience or motivation.

In Zhang et al.'s book-searching study, 96 participants were divided between two 48-participant substudies with and without gaze-assisted post-task feedback. The authors reported that the combined visual/audio benefit was conditional on task feedback [40]. This is relevant because assistance modality and feedback timing interact with task design. Their target book was specified to the user; the present task withholds identity and keeps the guidance mechanism common across voices. Their sample size is not a power justification for this experiment.

The current study does not include an unguided or gaze-unaware control. Both voice conditions receive the same gaze-contingent policy. Earlier project documents that considered a factorial combination of voice and gaze contingency are historical, and their interpretation cannot be transferred to this two-condition design. Any observed difference will concern the assigned voice experience within the shared assistance system.

Voice as a Social and Personalized Interface

Research on machine-generated speech has considered personality, source attribution, and social presence [29, 28, 23]. These perspectives motivate treating voice as more than a neutral transport for words. They do not establish that every listener will form the same impression, or that a social impression necessarily improves task performance.

For this project, personalization has two meanings. Catalog personalization allows the participant to choose among existing voices. Self-similar personalization uses recorded speech to create a synthetic voice resembling that participant. Both forms may produce familiarity or liking, but they operate through different routes. Preference is selected directly in the catalog condition, whereas perceived resemblance must be checked in the self-similar condition. The manipulation is therefore not adequately described as personalized versus generic.

The bibliography also includes work on voice-assistant personality customization and similarity [35]. Personality similarity is conceptually different from acoustic self-resemblance. A listener can prefer a speaking style without hearing it as their own voice, and a close acoustic resemblance can be experienced as uncomfortable. The planned manipulation check asks about perceived voice similarity, while the experience battery separately asks about liking, naturalness, familiarity, and uncanniness.

Self-Similar Voice and Embodied Collaboration

Kao et al. investigate self-similar avatar voices in educational games [19]. Guo et al. examine self-similar appearance and voice in a jigsaw-puzzle co-solving task with a virtual character [16]. These studies are closely related to the manipulation, but the surrounding interactions differ. An avatar can supply visual identity cues, educational games can include learning outcomes, and a collaborator can have a visible body. Follow My Voice uses a voice agent to guide spatial search without a visible agent body.

Guo et al. reported that self-similar voices increased likability and believability ratings while also increasing eeriness. Appearance self-similarity affected co-presence and perceived intelligence, and matched appearance/voice similarity influenced believability [16]. These results argue against treating resemblance as uniformly positive. They motivate retaining both favorable and unfavorable experience measures here. They do not predict faster object search, because their evidence concerns a different task and embodied-character configuration.

The distinction between embodiment of the participant and embodiment of the agent is important. Participants wear a headset and use a tracked controller, so the interaction involves physical sensing and action. The guiding agent, however, is not represented by a co-present animated character. Measures adapted from an embodied collaborator must therefore be described as adaptations. Their use does not establish that the same factor structure or interpretation holds for a disembodied voice.

Self-voice applications in daily reminders and emotional self-voice interventions extend the range of contexts represented in the bibliography [20, 13]. They motivate asking whether participants can imagine uses for their own synthetic voice. They do not establish that navigation guidance will be more effective or that participants will experience it as inner speech. The debrief conse-

quently begins with open questions before introducing probes about eeriness, self-talk, and possible applications.

Naturalness, Familiarity, and Uncanniness

Voice resemblance is not the only acoustic difference that listeners may notice. Synthesis can change articulation, prosody, loudness, duration, or perceived naturalness. Research on vocal uncanniness provides a reason to keep these qualities visible rather than subsume them under similarity [10]. A clone that is recognizable but difficult to understand is not equivalent to a recognizable and intelligible clone.

The experiment will use the same semantic library and provider for both conditions. That is an implementation control, not proof of equal acoustics. Prepared clips can differ in duration and pronunciation even when their text matches. Pilot checks will inspect intelligibility and comfortable playback levels across catalog and cloned voices. Actual provider, model, voice, and library identifiers will identify the configuration a participant heard.

Custom experience items and the appealing, eerie, and familiar items selected from Guo et al. will remain separate. Combining them would imply a measurement model not established for this study. A participant may report a voice as both familiar and eerie. Preserving the individual directions allows that response to remain interpretable.

Context and Short-Term Memory in XR Agents

EmBARDiment is a relevant IEEE VR example of an embodied agent using contextual information in XR [7]. The project draws on its context-management motivation while implementing its own

bounded history. The history here records recent spatial progress, actual speech, and guidance transitions. It does not provide long-term autobiographical memory or an inferred model of the participant's mental state.

Memory-Augmented AR Agents and GenLARP represent additional directions in the bibliography [9, 38]. The former is a conceptual framework and the latter a short system description; neither supplies a validated timing threshold for this application. User-aware virtual agents and proactive AR assistance also offer context for designing state-dependent behavior [15, 24]. These works position a design problem, rather than prove that this implementation's memory mechanism improves performance.

Short-term history is operationally useful when the same current geometry can have different meanings. A target to the left may be one the participant has not yet approached, or one they have just passed while turning right. A recently delivered cue can make another identical cue unnecessary. History distinguishes these cases only if it records reliable observations and actual playback. Holding this mechanism constant across voices means the study cannot estimate its independent effect.

Eye-Tracking Quality and Interpretation

Eye-tracking quality involves more than a nominal sampling frequency. The reporting guidance of Dunn et al. and the data-quality discussion of Niehorster et al. motivate reporting acquisition, validity, spatial performance, and missingness [12, 30]. The system must establish how many distinct source observations arrive, how old they are when used, and how error varies across the search space.

Validation work on the Vive Pro Eye and tools such as GazeMetrics are methodological resources [33, 34, 1]. They are not substitutes for validation on the HTC Vive Focus Vision.

Device model, fit, calibration, coordinate conversion, and task geometry can differ. The protocol therefore requires independent targets across heights, boundaries, distances, and head orientations. Validation targets must be separate from targets used to estimate an offset correction.

Drewes et al. compare gaze during locomotion in virtual reality and the real world [11]. This study is seated and uses head and chair rotation rather than free walking, so locomotion findings cannot be imported as direct estimates of its gaze behavior. Head motion will be logged because it participates in search, while claims about eye events remain bounded by the available measurements.

Salvucci and Goldberg discuss identifying fixations and saccades from eye-tracking protocols [32]. An interaction region around an object serves a different purpose. The 100-ms vicinity qualification in this system is an interface acknowledgment rule, not a validated fixation classifier. The thesis will report gaze contacts and qualified vicinity events under their own names.

Workload, Motivation, and Agent Perception

Zagermann et al. discuss eye tracking as a means of investigating cognitive load in visual computing [39]. That motivates considering gaze-derived measures, but does not license interpreting an unvalidated ray-hit count as workload. The current confirmatory plan uses subjective workload dimensions and task performance. Pupil, fixation, and saccade analyses remain contingent on separate acquisition and event-validation work.

NASA-TLX supplies six workload dimensions [17]. The project's 0–10 response range and unweighted complete mean are documented adaptations. The Intrinsic Motivation Inventory distinguishes interest/enjoyment from related subscales such as effort, pressure, and value [31, 26, 8]. Separate subscale reporting preserves that distinction. The instrument appendix specifies current wording and reverse keys so scoring can be reproduced.

Agent perception measures are drawn through Guo et al., with attribution to their originating measures [16, 6, 27]. Co-presence, perceived intelligence, and anthropomorphism concern participant attributions. A high rating does not establish sentience or decision correctness. Conversely, deterministic validity checks do not establish that a participant found the agent natural. The design needs both kinds of evidence while keeping their meanings separate.

Synthesis and Research Gap

The selected literature supports a study at the intersection of spatial assistance and voice identity. Gaze-contingent guidance provides a mechanism for adapting directions to search. Self-similar voice work motivates a personalized voice comparison. Eye-tracking and questionnaire methods establish requirements for measuring that comparison without overstating what the signals mean.

The specific gap addressed here is preferred versus self-similar voice delivery in an initially unnamed, surrounding object-search task under a common adaptive policy. The thesis does not claim that no related system exists. Instead, it makes the comparison operational: the task, shared mechanism, voice assignment, outcome boundaries, and measurement limitations are specified together. The answer will depend on the completed study rather than on an assumption that familiar voices are necessarily better.

CHAPTER 3: SYSTEM DESIGN AND IMPLEMENTATION

Follow My Voice is a gaze-contingent voice guidance system designed to help a seated user find a target object in a surrounding search scene. The system divides responsibility between a local guidance layer, which owns all spatial reasoning, and a cloud language model agent, Jev, which decides how and when to speak. Figure 3.3 summarizes the pipeline from sensing to speech output. This division makes spatial eligibility logic deterministic and traceable while allowing the agent to adapt the phrasing and timing of guidance to the user's recent behavior.

Apparatus

The application is developed in Unity 6 (6000.3.10f1) using the Unity OpenXR Plugin (1.16.1), HTC VIVE OpenXR integration (2.5.1), and Unity XR Interaction Toolkit (3.3.1). It is deployed as a standalone Android ARM64 application on the HTC VIVE Focus Vision head-mounted display, with rendering performed locally on the headset using the Universal Render Pipeline (17.3.0). The headset supplies head-pose and eye-gaze tracking, while the tracked right-hand controller supplies pointing and trigger input. Video passthrough supports the mixed-reality view. Network access supports cloud-based agent decisions and voice preparation; it is not used to stream rendered frames from a host computer.

Search Task

The search task takes place in mixed reality using video passthrough in a neutral physical room. Participants will sit on a swivel chair and interact using the right controller. Virtual colored geometric objects hang on eight vertical planes anchored along the surrounding walls. Figure 3.1

shows the layout from above. The application holds the target identity, which is withheld from the opening instructions, so participants must rely on voice guidance to locate it. Participants select an object by pointing the controller ray at it and pressing the trigger. Stimulus density and spatial bounds will be frozen for the study release after seated-comfort and ray-selection pilots. For each trial, the system logs object positions, layout seed, wall assignment, and target identity.

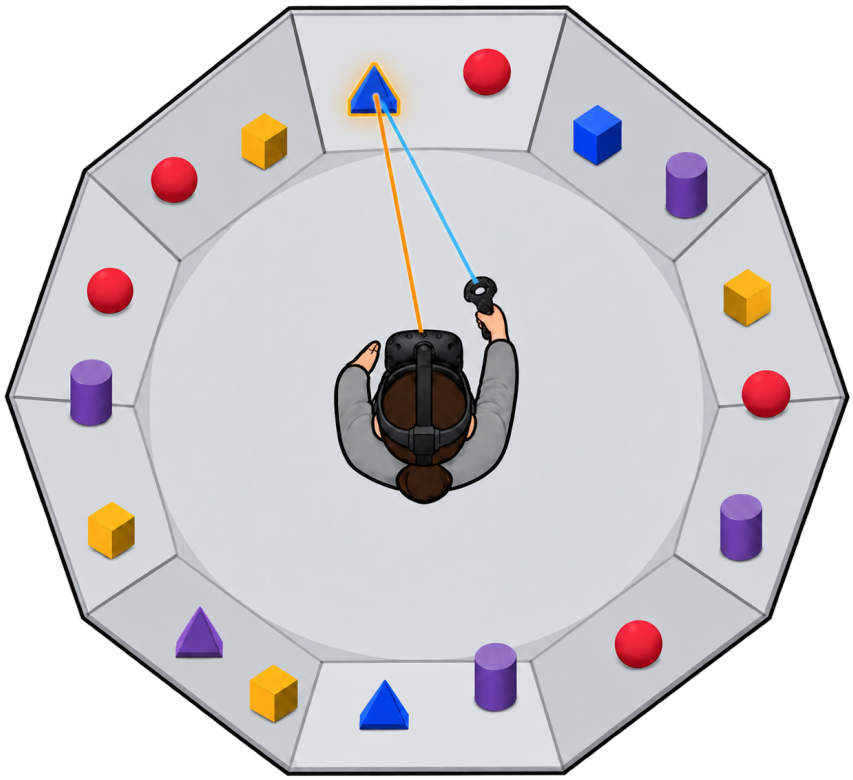


Figure 3.1: Top-down view of the search scene. The seated participant is surrounded by wall panels holding colored geometric objects. The gaze ray (orange) and controller ray (blue) converge on the highlighted target, the state in which the participant is ready to select with the trigger. The layout is illustrative; the drawn object count is schematic; object positions and target identity vary by trial. The converging rays and highlighted endpoint illustrate interaction, not an advance participant-visible target-identification cue.

Sensing and Guidance State

The system continuously collects eye-gaze origin, direction, and validity; head pose; right-controller pose; and ray-object intersections. A local state machine combines these observations with the target geometry to track the participant's progress through the search: turning toward the target wall, refining within the wall, acquiring the target with gaze, pointing with the controller, and becoming ready to select. Additional states cover suspended or paused operation. Each state transition is logged with its cause.

Because guidance is only as reliable as the evidence behind it, the state machine issues new gaze-dependent guidance only when gaze data are valid and recent. Gaze ray hits, gaze-vicinity candidates, and controller hits are treated as separate observations. When gaze-vicinity regions overlap, the closest object is chosen as the vicinity candidate. A controller ray hit on the target immediately suppresses contradictory directional cues. The spoken pointing acknowledgment separately requires a one-second controller hold with valid nearby gaze when speech is admitted. Suppression is intended to prevent away-directions during the qualification interval without falsely acknowledging an unqualified hold. Gaze confirmation, controller confirmation, and trigger selection are recorded as distinct events.

Short-Term Interaction Memory

To give guidance continuity, the system maintains two short-term memories. A spatial history tracks turning progress and supports cues that return the participant to the target. A speech-decision memory records what the system actually said and how the guidance state changed. Each speech entry stores its relative age, context, and whether playback completed or was interrupted; each state entry stores the prior and current states and the reason for the transition. The speech and transition lists

each retain at most four entries; snapshots sent to Jev exclude entries older than 15 seconds. This bounded recent context is related to the context-management motivation of EmBARDiment [7], but the current implementation does not use its object-admission or dropout policy. The separate spatial history retains up to three seconds or 384 target-relative gaze observations and clears on invalid evidence, source changes, or observation gaps longer than 100 ms. These limits are engineering settings, not validated fixation or blink thresholds. Gaze-vicinity confirmation currently requires 100 ms of valid qualifying evidence; invalid evidence resets its entry timer, and reacquisition after a spoken confirmation requires 150 ms of valid gaze away. Those interaction gates do not admit objects to speech memory, and missing gaze is not counted as observed dwell. Both memory types are cleared whenever guidance resets and is copied before each asynchronous request. This memory is deliberately scoped to the current task: it does not build a persistent participant profile or attempt to infer unobserved mental states.

Figure 3.2 shows which recent evidence accompanies an asynchronous request. Its time limits bound the software context and must not be interpreted as estimates of human short-term memory.

Agent Decisions and Playback Control

Jev selects from a constrained set of options rather than generating free-form instructions. Local geometry and state rules first filter the phrase library to candidates that are valid in the current situation. Each request to Jev contains a state snapshot, recent memory, playback status, and two to five eligible phrases. Jev then chooses the wording and whether to speak now or wait. Target definition and geometric reasoning remain local. Eligibility filtering and playback revalidation are intended to reduce spatially incorrect or stale guidance, but tracking error, imperfect geometry and software faults can still produce unhelpful instructions; recorded-sequence and headset checks remain necessary.

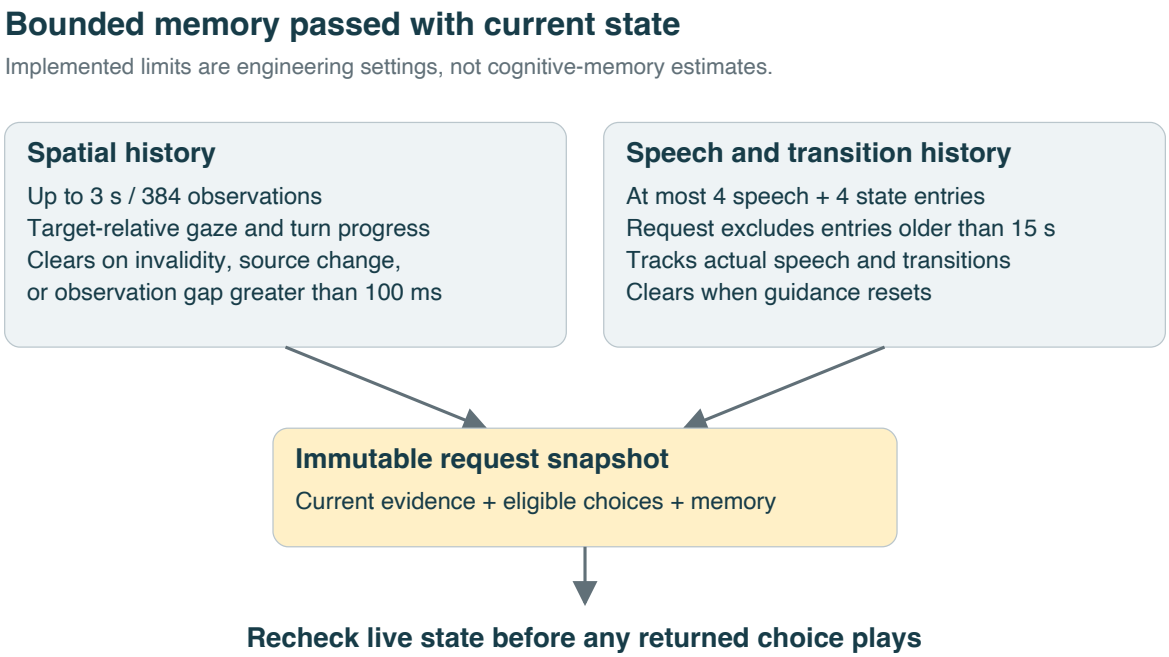


Figure 3.2: Implemented memory limits and snapshot boundary. Spatial history and speech/state history have different retention and reset rules. A copied request preserves its original context, while playback still requires a fresh state check.

Because cloud responses arrive asynchronously, the participant's state may change before a decision returns. A local speech scheduler therefore revalidates each decision against the current state, evidence, pacing, and object intersections before playback, and discards decisions that are stale or conflict with the current situation. Qualified gaze and controller acknowledgments and eligible turn-progress milestones can use local handling and are played from prepared phrases without waiting for a new cloud decision, subject to the same checks. Each completed or interrupted utterance updates both the memory and the event log. Model latency, deferral, cancellation, fallback, and playback onset are logged separately to support delay attribution. When a trace lacks the required events, the cause remains unknown rather than inferred from the spoken phrase alone.

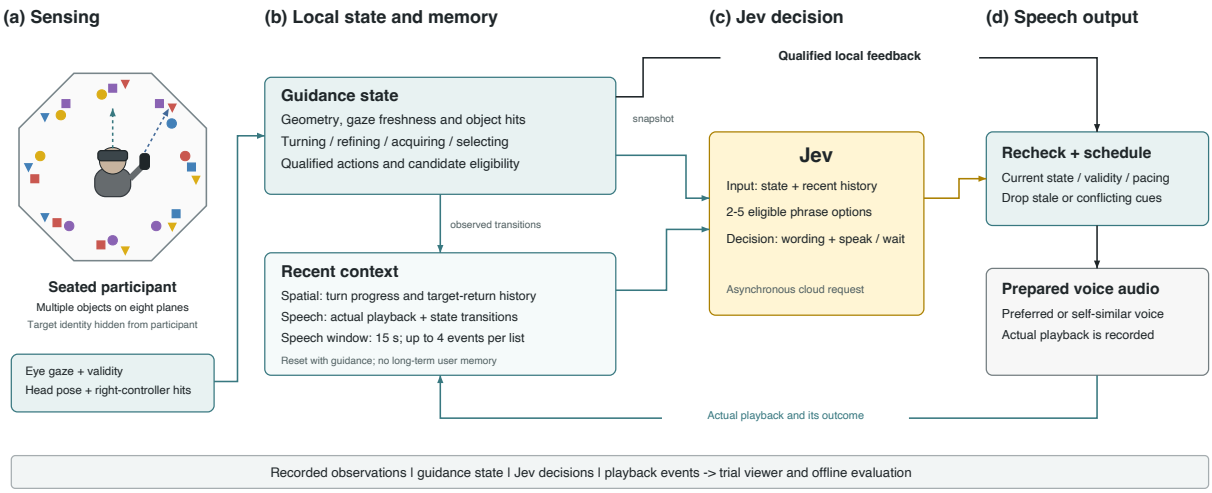


Figure 3.3: System architecture. (a) Gaze, head, and controller data are sensed. (b) A local state machine and short-term memory determine which guidance actions are valid. (c) Jev chooses wording and timing from eligible options. (d) The scheduler revalidates each cue against the current state before playback. Actual playback feeds back into memory. The scene inset is schematic and does not reflect stimulus count or any participant-visible target cue.

Voice Preparation and Telemetry

All voice audio is pre-synthesized with ElevenLabs, which supplies both the catalog voice and the self-similar voice. To support the self-similar condition, each participant records the same standard passage in their usual speaking voice. Both conditions share the semantic phrase library, guidance policy, and playback routing. Voice condition is the intended manipulation; behavior-dependent wording, hint counts and synthesized acoustics need not be identical. The preparation pipeline checks wording and library readiness. Per-voice intelligibility, loudness and acoustic comparability still require pilot verification; automated checks do not establish perceptual equivalence. No pitch or spatial effects are applied to the self-similar voice. We log the provider, model, voice ID, library version, and synthesis settings.

The Trial Review application aligns gaze and controller state, guidance transitions, agent decisions, and playback events with trial geometry on a common timeline. Trial-level aggregates are stored separately from the raw gaze, motion, and event streams. Effective gaze sampling rate must be verified from distinct source timestamps on the tested headset build rather than inferred from the nominal specification or application polling rate. Similarly, the analysis will report ray intersections as gaze hits and do not label them as fixations unless a validated event classifier is applied.

Trial Review Interface

Figure 3.4 shows the standalone Trial Review application with its bundled synthetic demonstration. The scene and inspector share a timeline with speech, gaze, events, and decision tracks. The application also supports receiving headset records for live review; transfer and receipt times are separate from source event times, and late arrival must not be interpreted as delayed headset guidance. This capture illustrates the review interface only; it is not a participant recording or a performance result. In particular, the demonstration explicitly marks state history as missing and highlights as reconstructed. That distinction is useful when reviewing legacy records: the viewer should expose unavailable evidence rather than imply it was measured.

Ownership and State Transitions

The central guidance machine is a plain C# object, while Unity MonoBehaviours connect it to tracked devices, scene objects, and audio. This separates deterministic state transitions from engine callbacks. The round manager owns readiness, trial onset, completion, pause intervals, and block progression. The guidance machine owns the interpretation of evidence during search. The speech



Figure 3.4: Actual application screenshot captured for this draft using the bundled SYNTHETIC / Run 1 demonstration. The central reconstruction, right-hand inspector, and timeline provide synchronized review. Missing state, estimated rings, and reconstructed highlights are labeled by the application. This is a visualization example, not headset video or empirical study data.

scheduler owns pending requests, cancellation, and playback admission. These are coordinated responsibilities rather than a single global state variable controlling the whole application.

A state label is useful only when it is accompanied by its evidence. A refining state indicates that guidance has entered the target wall, but does not prove that the eyes intersect the target object. An acquired state can reflect qualified vicinity evidence, but does not mean a trigger selection occurred. The viewer therefore presents state transitions alongside gaze-hit, vicinity-candidate, controller-hit, and playback records. This makes it possible to inspect an apparent disagreement rather than assume the label is complete ground truth.

Table 3.1: Guidance responsibilities and evidence boundaries.

Component	Owns	Does not establish
Observation adapters	Validity, poses, ray hits, source timing	Recognition or fixation
Guidance machine	Stage, eligible action, bounded history	Correct trigger selection or conscious recognition
Jev decision provider	Choice among eligible phrases; speak or wait	New target identity or unchecked geometry
Speech scheduler	Current-state check, pacing, cancellation, start	Acoustic arrival at the ear
Round manager	Readiness, reveal, selection, pauses and blocks	Independent validation of gaze accuracy
Trial Review	Reconstruction from saved evidence	Missing observations between recorded samples

The usual path is turning toward a wall, refining within it, acquiring the target with gaze, and preparing to select. This is not a mandatory one-way chain. Participants can overshoot, revisit a wall, point before gaze qualifies, or lose tracking. The implemented state families represent these returns; release testing must verify them without repeatedly announcing a new arrival whenever the participant looks upward on the same wall. Invalid or stale evidence suspends new gaze-dependent decisions rather than treating the absence as a departure.

Figure 3.5 summarizes the state families. The arrows show a common progression rather than compulsory transitions: correct right-trigger selection can complete an active trial without a preceding spoken gaze or pointing acknowledgment.

A Worked Guidance Sequence

Consider an illustrative target near the wall opposite the starting cross, with a rightward offset or the documented backward-turn tie rule. The opening cue may say that the object is to the right while objects remain hidden. Once search begins, progress cues can acknowledge continued rightward

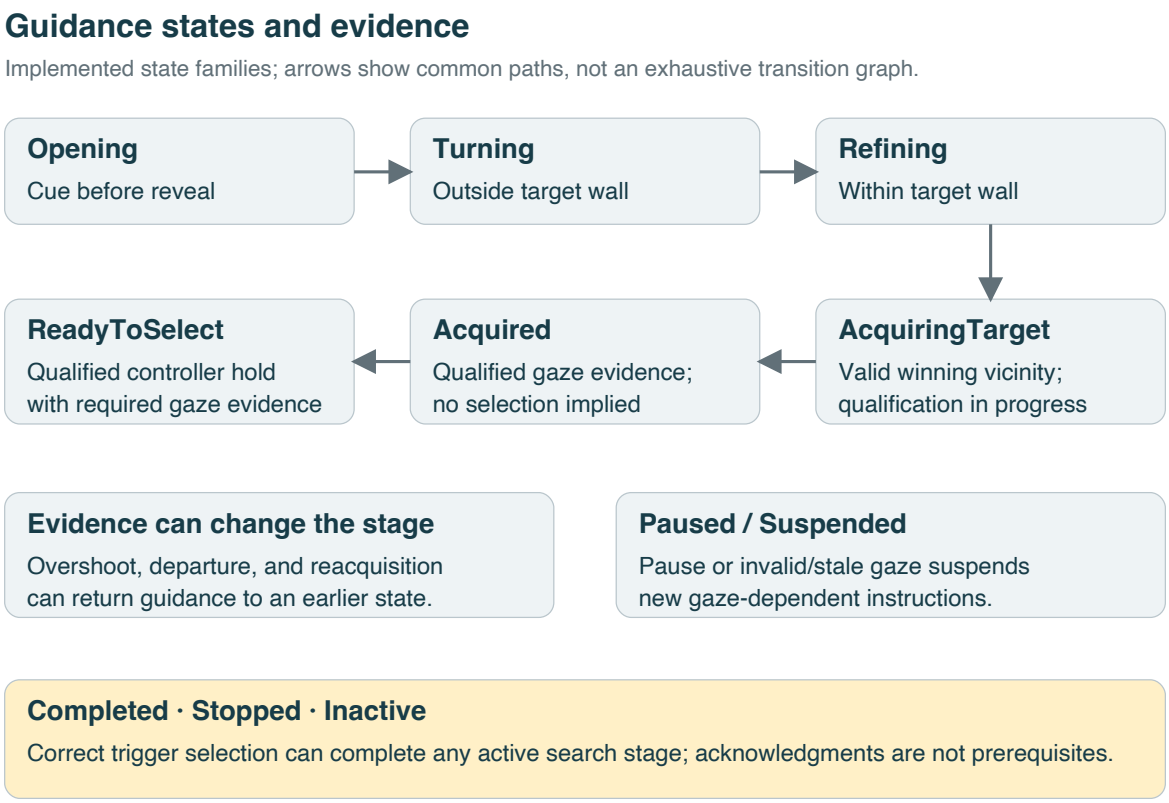


Figure 3.5: Implemented guidance state families. Turning, refining, acquisition, and readiness are evidence-dependent interpretations. Tracking loss, pauses, target departure, and reacquisition can change the path. Trial completion remains owned by the round manager.

turning. On stable entry to the target wall, assistance can move to within-wall refinement, such as looking lower and slightly left. The reference for that refinement must be the recorded gaze geometry and local target relation, not a neighboring object's apparent screen position.

If the participant passes the target, recent spatial history can support a return cue. If the participant instead looks near the target long enough, a gaze acknowledgment becomes eligible. If the controller is already intersecting the target, away-directions are suppressed immediately. A qualified pointing acknowledgment can then identify controller evidence, and a right-trigger press can

complete the trial. This example illustrates intended transitions; it is not a measured trajectory or evidence that all people follow this sequence.

Asynchronous Decisions and Cancellation

Each cloud request receives a snapshot, not a live reference to mutable state. The snapshot includes current evidence, relevant geometry, speech status, and bounded recent history. A response can become stale because the participant moves, tracking changes, the voice changes, a pause begins, or the trial ends. The scheduler checks those conditions again before playing the chosen clip. A valid answer to an old question is not necessarily a valid instruction now.

Prepared audio removes synthesis from the immediate hint path, but does not remove network latency from cloud choice. Time-sensitive qualified feedback can use local admission without waiting for another Jev decision. Ordinary hints can wait for phrasing choice within bounded scheduling rules. Interruption policy balances urgency against comprehensibility; finishing a short useful phrase may be preferable to cutting it at an awkward boundary. The documented implementation uses a 150-ms fade as a playback technique, not a guarantee of natural conversation.

Experimental policies that allow Jev broader action choice must be versioned separately from the participant policy. Replaying the same observations through multiple policies can compare decision latency and admissibility. It cannot show how a participant would have moved after hearing a different cue. The participant release will need one frozen policy and a manifest of thresholds, phrase library, Jev provider/model identifier, version and available sampling settings. Whether the provider sees voice-condition identifiers must be verified against the frozen payload; this draft does not assume condition-blind decisions. That audit will distinguish voice-related cancellation keys from information actually transmitted to the provider.

CHAPTER 4: EXPERIMENTAL METHODOLOGY

Draft status. This chapter specifies the proposed confirmatory procedure. The implemented snapshot supports one practice, two 40-trial blocks, and final questionnaire handoff. Exact pre-cue geometry, vertical balancing, sample size, and pilot acceptance rules remain unresolved, as detailed in Appendix D.

Task

Participants will perform an agent-guided object-search task while seated on a swivel chair. Multiple colored geometric objects will be scattered across eight surrounding vertical planes. The application will know the target identity, but the participant will not receive a color-and-shape description before searching. Instead, the agent will guide the participant toward the target wall and then refine the direction horizontally, vertically or diagonally. The participant will select an object using the right-controller ray and trigger. Looking at an object, receiving a gaze confirmation or pointing at it will not itself complete the trial.

Each trial will begin with the participant looking at a readiness cross. After readiness is established, an opening directional cue will play while the objects remain hidden. Object reveal will start active search timing. A correct trigger selection will stop timing and produce success feedback; an incorrect selection will be counted and search will continue. This separates readiness and opening guidance from the timed search interval while retaining their timestamps for analysis.

Design and trial allocation

The study will use a within-participant comparison of two voice conditions: a participant-preferred catalog voice and a self-similar synthetic voice. Both conditions will use the same guidance policy, semantic phrase library, selection method and visual-feedback settings. Actual phrase sequences may differ because guidance responds to behavior. The comparison will therefore concern two voice configurations within an adaptive assistance system; it will not estimate the effect of gaze contingency itself. The catalog voice is participant-chosen while the clone is derived from the participant. Preference, gender presentation, and synthesis quality may consequently differ along with self-similarity; the contrast cannot isolate self-similarity from all of these features.

A saved, balanced randomized allocation will assign preferred-first or self-similar-first order. Each participant will complete two measured blocks of 40 trials, one per voice, following one practice trial. The practice trial will use the first voice in the assigned order: preferred catalog voice for Order A, or self-similar voice for Order B. Practice will be excluded from measured outcomes. A seed will make the trial and layout schedules reproducible. Nominal horizontal-angle quotas will be identical across blocks; trial order and left/right assignments will be randomized subject to those quotas. The implemented schedule samples left/right signs without enforcing an equal split within each bilateral angle category. Exact four-left/four-right balance is therefore not guaranteed; adopting that additional constraint would require an explicit protocol and implementation change before freezing the study.

Objects will vary in within-plane position, so the exact starting angle and target height will be recorded in addition to the assigned wall. The vertical allocation will use target-center height relative to the readiness cross, not the participant's seated eye height. Feasible height ranges and realized quotas will be checked in pilot schedules before the study release is frozen.

Table 4.1: Study allocation. Nominal wall angles are schedule categories, not exact object bearings or accumulated head rotation.

Parameter	Allocation
Voice conditions	Preferred catalog voice; self-similar synthetic voice
Measured trials	40 per block; 80 per participant
Practice	One trial; excluded from outcomes
Absolute nominal wall step	Eight trials each at 0, 45, 90, 135 and 180 degrees per block
Vertical allocation	20 targets above and 20 below readiness-cross height per block (proposed; enforcement remains an implementation requirement)
Condition order	Counterbalanced preferred-first and self-similar-first
Selection	Right-controller ray and a trigger press

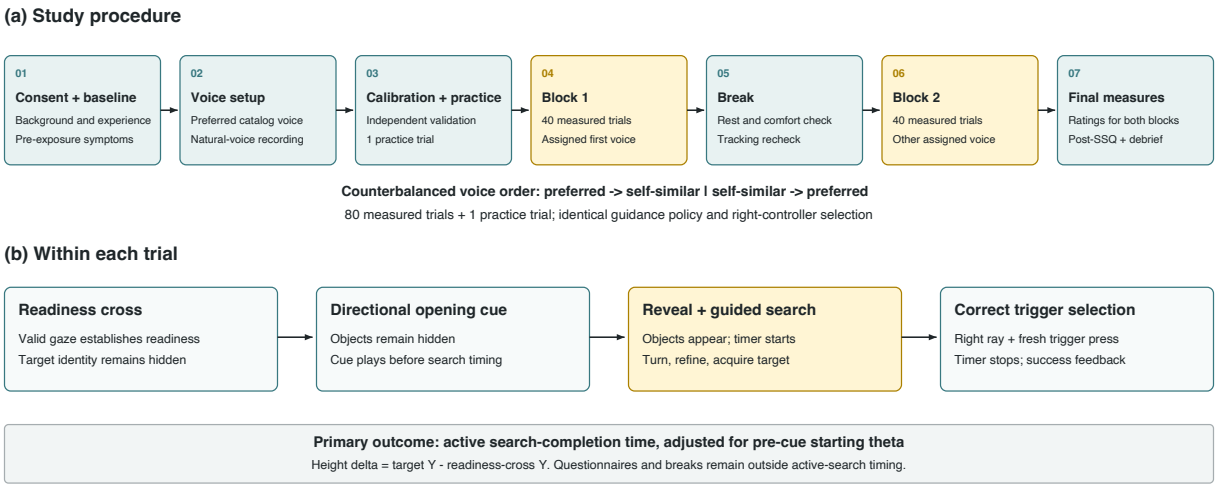


Figure 4.1: Study procedure and trial timing. Opening guidance precedes object reveal; active completion time ends at correct right-trigger selection. Questionnaires and between-trial breaks remain outside active-search timing.

Nominal angle steps are relative to the saved readiness/start wall, represented in the aligned schedule as the previous target wall center. The first measured start wall uses the practice target wall; the completed-session check must verify the saved cross placement. This wall-center relationship is distinct from the proposed exact cross-to-target bearing.

Measures

Task performance

Active completion time.

The primary outcome will be elapsed active-search time in seconds from object reveal (t_r) to the correct right-trigger selection (t_s), excluding the elapsed duration of recorded pauses overlapping that interval:

$$T_{\text{active}} = (t_s - t_r) - \mu([t_r, t_s] \cap \mathcal{P}) .$$

Here $\mathcal{P}$ is the union of valid recorded pause intervals, and μ gives elapsed duration. Lower values will indicate faster completion. Readiness, opening guidance, post-selection feedback, questionnaires and between-trial breaks will not contribute to this outcome. An unfinished trial will retain its observed exposure and outcome, but will not receive a successful completion time. The full interval from opening-cue onset to selection will be retained separately to describe the complete instructed task duration.

Target contact and selection errors.

Time to first target gaze will be the active-clock time of the first valid, recorded gaze-ray intersection with the target. Time to first controller contact will use the first tracked right-controller ray–target intersection. Last target-gaze contact time will denote the last observed valid target intersection during the trial, not a fixation duration or a spoken confirmation. These event times will remain distinct from trigger selection and from the larger gaze-vicinity region used by the guidance system. Incorrect selections will be counted as trigger selections of non-target objects. If no valid target contact is observed, its latency will remain missing rather than being set to zero. Contact times will retain their source and sampling resolution.

Starting geometry and search behavior

Starting angle, height and distance.

Before the opening cue, the system will record the reference origin, readiness-cross position and target position. Starting theta (θ_0) will be the signed angle about world up between the horizontal origin-to-cross and origin-to-target vectors. Its magnitude $|\theta_0|$ will quantify horizontal turn demand, and its sign will preserve direction. Degenerate horizontal vectors will be marked unavailable. The primary analysis will adjust for $|\theta_0|$ rather than divide completion time by degrees. Nominal wall-angle assignment, exact pre-cue theta, and head-to-target error observed after reveal will remain separate variables. The exact pre-cue snapshot is not yet implemented in the aligned study build. Existing nominal-theta fields are not substitutes; this snapshot must be implemented and verified before the study release is frozen.

Signed target-height difference will be $\Delta h = y_{\text{target}} - y_{\text{cross}}$, in metres; positive values will indicate above-cross targets. Euclidean target distance will be the three-dimensional distance between target center and the first valid tracked head position after reveal, with that observation's timestamp retained. This post-reveal distance will not substitute for a missing pre-cue geometry snapshot. All positions used in a calculation will share a recorded coordinate frame.

Head movement and zone coverage.

Gross head translation will sum distances between adjacent valid head positions, in metres. Gross head rotation will sum shortest quaternion angular displacements between adjacent valid orientations, in degrees. Fixed-axis component totals will be retained separately and will not be mislabeled as Euler yaw, pitch and roll. Invalid intervals, pauses and excessive sample gaps will break accumulation rather than be interpolated across; covered time and gap counts will accompany totals. These measures describe head movement, not independently tracked chair rotation.

Distinct walls visited, wall transitions and revisits will summarize search coverage using stable valid gaze observations and a versioned zone-assignment rule. Transitions across tracking gaps will not be inferred. More motion or more visits will not by themselves establish worse strategy, since geometry and available observations also affect these totals.

Instructions and guidance timing.

Instruction count will include guidance utterances whose playback actually starts during active search, including delivered repeats. Requested repeats, delivered repeats and interrupted utterances will be retained separately; model requests, prefetched audio and success speech will not be counted as delivered search instructions. Opening cues and success feedback will retain separate event

labels. Timing analysis will distinguish evidence qualification, decision-request latency, model deferral, scheduling delay and audio onset. Gaze confirmations and controller confirmations will be evaluated separately against their corresponding gates; the controller gate includes both its hold and the required nearby valid gaze.

Subjective measures

Workload.

Six NASA-TLX dimensions—mental, physical and temporal demand, performance, effort and frustration—will be rated separately for each voice block during the final survey, after both blocks. The study will use a documented 0–10 adaptation, without pairwise weighting. Performance will run from Perfect to Failure; other dimensions will run from Very low to Very high. Individual ratings and the unweighted mean of all six completed dimensions will be retained; higher scores will indicate greater workload or poorer perceived performance. The source instrument is Hart and Staveland's NASA-TLX.

Motivation and activity experience.

The Intrinsic Motivation Inventory will contribute interest/enjoyment (7 items), effort/importance (5), pressure/tension (5), and value/usefulness (7). Responses will range from 1 = Not at all true to 7 = Very true. Published reverse keys will be applied before averaging each complete subscale. Interest/enjoyment will represent intrinsic motivation; the four subscales will not be combined into an overall IMI score. Activity wording will refer to the completed search block, with task-specific adaptations documented.

Agent perception.

Twenty-one items from Guo et al. (2024), Appendix A, Table A1 will assess co-presence (Q1–6), perceived intelligence (Q13–18), appealing/eerie/familiar impressions (Q20–22), and perceived anthropomorphism (Q37–42). "Virtual character" will be adapted to "voice agent." Responses will range from 1 = Very strongly disagree to 7 = Very strongly agree. The three six-item constructs will have separate complete-scale means. Q20–22 will be analyzed individually, without a combined eeriness score. These attributed qualities will not be interpreted as actual agent intelligence or sentience.

Voice experience and reliance.

Twelve custom 1–7 agreement items will assess voice similarity, responsiveness, helpfulness, distraction, familiarity, trust, comfort, uncanniness, reliance, naturalness, willingness to follow and motivation to act without guidance. Voice similarity will serve as a manipulation check. Voice liking and retrospective percentage of guidance use will each use 0–100 responses with separate unable-to-answer options. These 14 numeric items will be analyzed separately; they are not a validated composite scale. Liking during voice audition will remain separate from condition-specific liking collected at the end.

Symptoms and debrief.

The 16 Simulator Sickness Questionnaire symptoms will be rated from 0 = None to 3 = Severe before headset exposure and after headset removal. Canonical total or subscale scores will require the published membership and weights. Open debrief questions will explore perceived performance,

guidance experience, own-voice comfort, possible use cases and voice preference; broad experience questions will precede specific eeriness or inner-dialogue probes.

Each condition-specific set in the final survey will therefore contain 65 numeric responses (130 across both sets): six workload ratings, 24 IMI items, 21 Guo items and 14 custom ratings. Technical checks and optional text will be additional. Questionnaire observations will have participant-by-block scope even when copied into per-trial summary files; they will not be treated as 40 independent responses. Missing answers will not be replaced with zero.

Table 4.2 summarizes the condition-specific battery. Counts refer to numeric items in one block-referenced set, not independent scales. Exact wording and reverse-key positions are preserved in Appendices A and B.

Table 4.2: Proposed final questionnaire battery for each referenced voice block. These ratings are administered after both blocks; setup voice ratings and the baseline/final SSQ are separate.

Measure	Items	Scoring and interpretation
NASA-TLX adaptation	6	Six 0–10 dimensions; complete unweighted mean. Preserve explicit bounds and adverse performance direction.
IMI	24	Separate interest/enjoyment (7), effort/importance (5), pressure/tension (5), and value/usefulness (7) means. No overall total.
Agent perception	21	Three six-item means: co-presence, intelligence, anthropomorphism. Appealing, eerie, and familiar remain three single items.
Custom voice and assistance	14	Individual ratings of voice and guidance experience. No unvalidated combined score.
Total per block reference	65	130 numeric ratings across two references, before technical checks and optional text.

Measurement quality

Tracking validity, observation age, data coverage, gaps and spatial-validation error will accompany gaze-derived measures. Ray intersections and vicinity hits will not be labeled fixations, and render-frame observations will not establish native eye-tracker sampling frequency. Pupil-based workload inference and fixation/saccade outcomes will require independent validation before inclusion. Trial performance and subjective summaries will remain linked to separate raw gaze, motion and event records through participant, session, block and trial identifiers. Scale metadata will distinguish the present 0–10 workload form from any earlier exports using another range.

Research questions

The primary question will be whether self-similar and preferred catalog voices differ in active completion time after accounting for absolute starting theta and repeated observations within participants. Secondary questions will concern workload and the four IMI subscales. Exploratory questions will concern agent perception, eeriness, voice experience, reliance, search behavior and variation in the voice contrast with starting angle. The primary voice contrast will be two-sided; a directional benefit of self-similarity will not be assumed. Repeated trials will not be treated as independent participants.

Participants

Participants will be consenting adults able to understand the task, hear the guidance, view the scene with usual correction and safely use the seated setup and right controller. Handedness, visual correction and prior XR and voice-assistant experience will be recorded. Recruitment size and the en-

rollment ceiling will be determined before recruitment using power analysis for the theta-adjusted repeated-measures outcome, including plausible within-participant variation and incomplete observations.

Procedure

Before participant use, consent materials must be reconciled with ElevenLabs voice processing and Jev cloud decisions, including the data transmitted, retention, access, and deletion arrangements. That reconciliation is a readiness requirement rather than a claim of current approval. After consent, participants will complete background questions and baseline symptom ratings. They will choose a preferred catalog voice, rate auditioned voices, record a standard passage in their natural voice for synthesis, and confirm comfortable, clear audio for each condition. The retained English recording passage is reproduced in Appendix A; its match to the frozen build and typical recording duration remain to be verified. Preferred gender presentation will be chosen explicitly rather than inferred from participant gender.

The researcher will fit the headset and conduct vendor calibration followed by independent spatial validation across relevant heights, wall boundaries, distances and head orientations. If an offset correction is introduced, its fitting targets will be separate from independent validation targets; this is a conditional validation requirement, not a claim that a correction is already validated. Standardized instructions will explain that participants will rely on the agent to discover the target and use the right trigger to select it. Participants will then complete the single practice trial and the two assigned 40-trial blocks.

No questionnaire will interrupt the block sequence. During the final survey, the same condition-specific questionnaire will be completed twice, once referring to Block 1 and once to Block 2, with

identical item wording and order. Breaks will be offered between trials and at the block boundary; any necessary in-trial pause will be logged immediately. Tracking will be rechecked between blocks and after headset adjustment. After headset removal, participants will first complete the final symptom check, then the two retrospective condition-specific sets, comparative voice preference and debrief. Participants may request a pause or stop, and symptoms or technical problems will prompt researcher assessment.

Figure 4.2 makes the distinction between administration time and reference scope explicit. Both condition-specific batteries occur in the final session; each response still refers to one measured block.

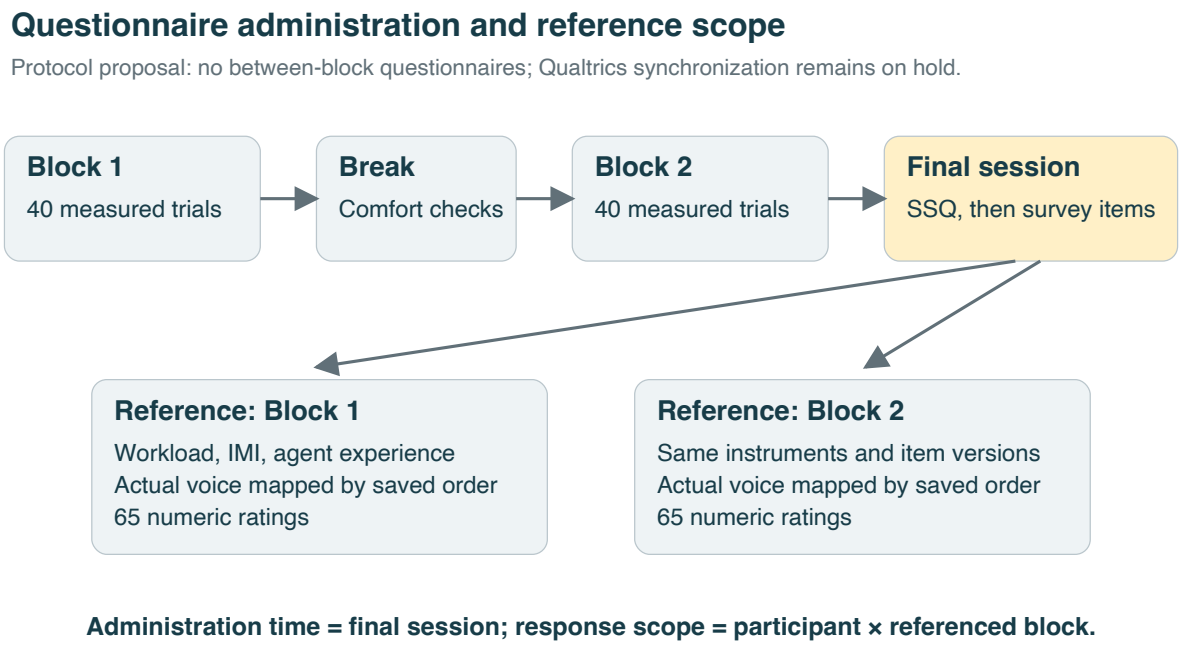


Figure 4.2: Protocol proposal for final retrospective administration. The intervening break contains comfort and tracking checks, not a block questionnaire. Each block-referenced battery contains 65 numeric ratings; the final SSQ is separate. Actual voice is recovered from the saved counterbalanced order.

The final survey will need an unambiguous reminder linking each block reference to its experienced voice. The reminder format (such as neutral labels or replayed samples) remains an open protocol decision and must be piloted without adding differential cues between conditions.

Pilot testing and parameter checks

Before confirmatory recruitment, pilot sessions will verify the realized angle/height schedules, target visibility and seated comfort, readiness and selection behavior, gaze validity, audio timing and wording, and reconstruction of trial events from the saved data. Survey comprehension and completion time will also be measured. Calibration acceptance, practice proficiency, timeout and technical-stop criteria will be set from those checks and recorded with the study build and questionnaire versions. Pilot observations will be reported as such rather than presented as study outcomes or evidence for a voice-condition effect.

CHAPTER 5: MEASUREMENT AND DATA PROVENANCE

From Interaction Events to Research Variables

The recorded data serve two purposes: reconstructing what happened in a trial and estimating the study outcomes. Reconstruction requires observations and events at their original temporal resolution. Statistical analysis usually requires a more compact table. The project therefore keeps trial summaries separate from gaze, head/controller motion, object layouts, and event streams. The summary is a derived representation, not a replacement for the underlying evidence.

Every derived value needs an identifiable source and a definition. A number called "time to gaze" is ambiguous unless it specifies the clock origin, the contact rule, validity handling, and the treatment of missing observations. The same is true of an angle: wall assignment, gaze error, starting geometry, and accumulated rotation can all be expressed in degrees, but describe different quantities. The data contract preserves these meanings in separate fields.

Figure 5.1 shows why a compact trial CSV must retain links to the underlying observations. Statistical scope and reconstruction detail impose different requirements on the exported data.

Coordinate Frames and Starting Geometry

Let $\mathbf{o}$ denote the saved trial reference origin, $\mathbf{c}$ the readiness-cross center, and $\mathbf{p}$ the target center. All three positions must be expressed in the same world coordinate frame. Let $\Pi(\mathbf{v}) = (v_x, 0, v_z)$ project a vector onto the horizontal plane. Define

$$\mathbf{a} = \Pi(\mathbf{c} - \mathbf{o}), \qquad \mathbf{b} = \Pi(\mathbf{p} - \mathbf{o}). \tag{5.1}$$

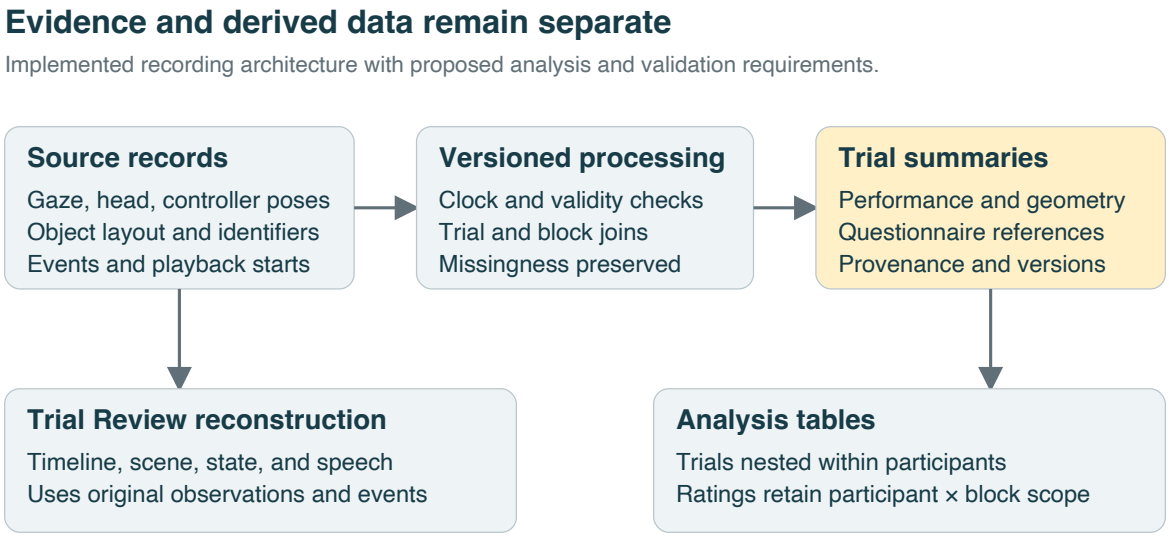


Figure 5.1: Recording and processing relationships. Source streams support reconstruction; derived summaries support analysis. Processing and scoring require versioned definitions. Repeating block-level ratings in trial rows does not create new independent observations.

For nondegenerate vectors, the exact pre-cue starting angle is

$$\theta_0 = \operatorname{atan2}(\hat{\mathbf{y}} \cdot (\mathbf{a} \times \mathbf{b}),\ \mathbf{a} \cdot \mathbf{b})\ \frac{180}{\pi}. \tag{5.2}$$

The sign convention will be verified against Unity's world-up signed-angle convention with known left and right targets. Its absolute value, $|\theta_0|$, describes horizontal demand from the saved cross reference. Exactly backward targets require a documented tie convention, because -180 and $+180$ identify the same horizontal bearing. The absolute magnitude remains 180 degrees in either case.

This calculation is a required addition to the frozen study data contract. The October 20, 2026 implementation alignment audit (source register in Appendix D) found that existing signed and

absolute theta fields represented nominal wall geometry relative to the previous target wall. They must not be renamed or silently interpreted as exact pre-cue measurements. A new snapshot must save origin, cross, target, timestamp, frame, and coordinate-frame identifier before the opening cue. If a horizontal vector is degenerate, the angle is unavailable; substituting zero would falsely describe an easy starting configuration.

Nominal categories remain useful for allocation. Each block contains eight scheduled trials in each absolute category, $0, 45, 90, 135, 180$ degrees. An object can be positioned away from the center of its assigned wall, so its exact bearing need not equal the nominal category. Balance in categories therefore does not prove exact geometric matching between voice blocks. Both assigned and realized geometry will be inspected.

The signed height difference is

$$\Delta h = p_y - c_y. \tag{5.3}$$

Positive values identify targets above the cross center and negative values identify targets below it. This reference is not interchangeable with seated eye height. The planned 20-above/20-below allocation must be checked against realized positions and a prespecified near-zero rule. At the source snapshot, floor-to-height scattering did not enforce that balance.

The three-dimensional distance measure uses a different observation: the first valid tracked head position $\mathbf{h}_k$ after reveal. Its value is $d = \|\mathbf{p} - \mathbf{h}_k\|_2$, in metres. The timestamp of $\mathbf{h}_k$ must accompany the distance. This post-reveal measurement describes geometry at an observed head pose; it cannot recover an unrecorded pre-cue snapshot.

Figures 5.2 and 5.3 illustrate the proposed geometric contract. Neither diagram represents a recorded layout. Their purpose is to show why nominal wall assignment, exact cross-relative bearing, cross-relative height, and live gaze error require separate names and saved references.

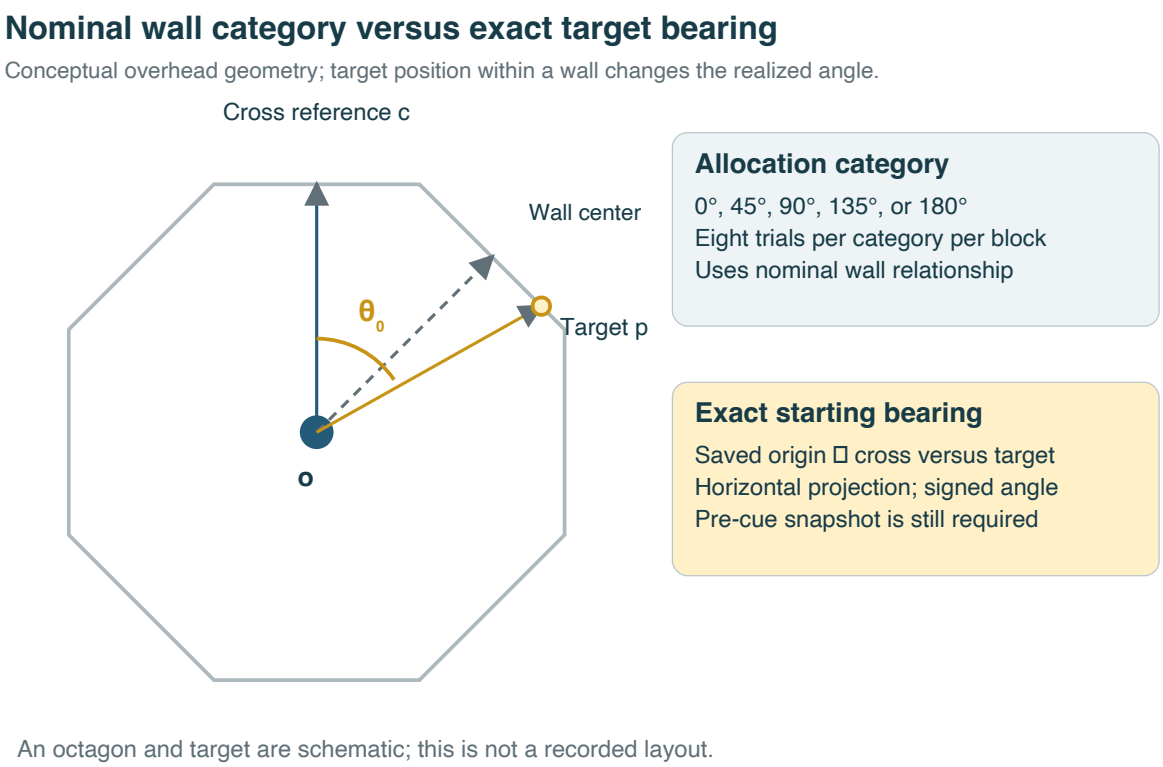


Figure 5.2: Conceptual overhead geometry. The target can be offset from its wall center, so an allocation category need not equal the realized starting bearing. The exact pre-cue snapshot remains an implementation requirement at the aligned checkpoint.

Trial Clocks and Pause Intervals

Let t_r be object reveal, t_s correct selection, and $\mathcal{P}$ the union of valid pause intervals. Active completion time is

$$T_{\text{active}} = (t_s - t_r) - \mu\big([t_r, t_s] \cap \mathcal{P}\big)\,, \tag{5.4}$$

where μ denotes elapsed duration. Using the union prevents overlapping pause reasons from being subtracted twice. For example, an application-focus pause can overlap a researcher-requested pause. The resulting active interval should remain stopped until all relevant pause reasons are resolved.

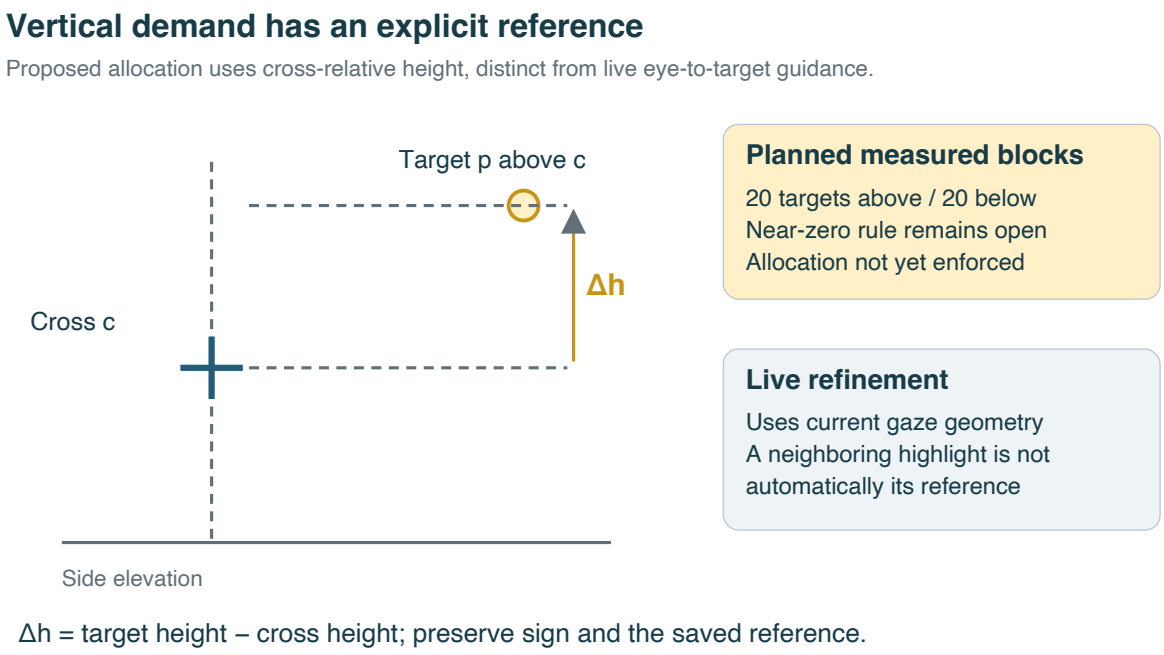


Figure 5.3: Conceptual side elevation for the proposed signed height difference. The target shown above the cross has positive height delta. The planned above/below quota and near-zero rule remain unresolved implementation and protocol checks; live vertical guidance uses current gaze geometry.

Readiness and opening speech precede reveal and are excluded from active search. This is an intentional outcome boundary, not missing task time. The system will retain a broader opening-cue-to-selection interval so that analyses can describe the time spent receiving the initial instruction. Prepared clips can have different durations across voices, making that additional record useful even when active completion remains the primary outcome.

A stopped or timed-out trial has an observed duration and an outcome reason, but does not have a successful completion time. The analysis must not copy its stop timestamp into a success field. Similarly, a late-arriving file in the viewer is a transfer delay, not necessarily a delay in the headset event. Source timestamps and receive timestamps have distinct roles.

Gaze Contact, Vicinity, Pointing, and Selection

Four event families are central to interpretation. A valid eye ray can intersect the target collider. A gaze direction can satisfy the larger target-vicinity rule. The right-controller ray can intersect the target. Finally, a trigger press can select the target. These events can occur at different times and need not occur in that order in every recorded sample sequence.

The first gaze-contact latency uses the first valid recorded target-ray intersection after reveal. The last gaze-contact time uses the last such intersection before trial end. Neither is a fixation duration. If no contact is observed because of a blink, acquisition gap, or a path that only enters the broader vicinity, the contact field remains missing. The presence of a spoken acknowledgment does not justify filling in an unobserved ray contact.

Vicinity qualification is an interaction rule. In the documented system, a target can be acknowledged after 100 ms of valid qualifying evidence, with overlapping regions resolved to a single closest candidate. The rule trades exact collider intersection for a tolerance intended to make feedback usable. That tolerance must be logged and piloted. The study-release tolerance remains to be frozen; the synthetic viewer screenshot's default ring is not a validated study setting. A very large region can create premature acknowledgment even if its timing is correct.

Controller pointing is a distinct input observation. Its spoken acknowledgment is not independent of gaze because the admission rule also requires valid nearby gaze. The documented pointing acknowledgment requires a qualified one-second controller hold with valid gaze in the target vicinity when speech is admitted. A valid target hit suppresses contradictory away-directions before the hold completes. This suppression prevents a direction such as "go back down" from being delivered while the controller is already on the target, without falsely saying the hold has qualified. Correct

trigger selection is the only event that successfully completes the search; a stop or timeout has a different outcome.

Figure 5.4 contrasts the three feedback and completion events. The example utterances communicate their meanings and are not asserted to be verbatim clips in every voice library.

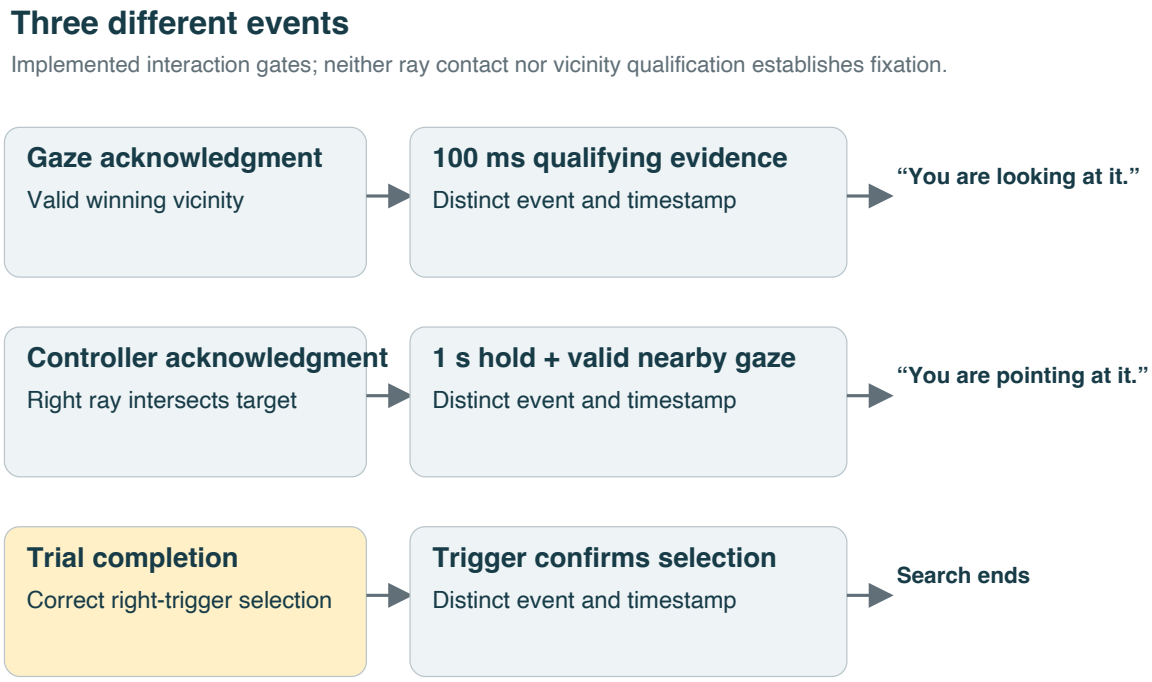


Figure 5.4: Implemented interaction distinctions. Gaze acknowledgment uses a 100-ms qualifying vicinity gate; pointing acknowledgment uses a one-second controller hold with required valid nearby gaze. A correct trigger selection completes the trial independently of whether either acknowledgment was spoken. These are interaction gates, not validated fixation measures.

Acquisition Rate, Validity, and Gaps

A render loop can query an eye-tracking API frequently without receiving a new hardware observation each time. The effective acquisition rate must therefore be estimated from distinct source

observations. If n unique observations span source times t_1 through t_n, a basic rate summary is $(n-1)/(t_n - t_1)$. The distribution of successive intervals, repeated timestamps, invalid observations, and long gaps must accompany that average. A mean alone can conceal bursty delivery.

Application timestamps describe when data were handled. Source timestamps describe when the tracking observation was produced, if the API exposes them reliably. Clock-domain mapping must be established before computing age or latency across these clocks. If unique source timing is unavailable, the report will state that limitation rather than infer 120-Hz binocular delivery from polling or headset specifications.

Spatial validation will compare gaze direction with known independent targets. Angular error will be computed from normalized gaze and expected target directions, with invalid observations excluded from the error distribution but retained in availability summaries. Precision and bias should be considered separately. A stable offset and a noisy signal can produce different guidance failures even if their average angular error is similar. The validation procedure must also describe target exposure, head pose, corrective lenses, and headset adjustments.

Motion, Coverage, and Instructions

Gross translation will sum Euclidean differences between adjacent valid head positions. Gross rotation will sum shortest quaternion angular differences. Invalid intervals, excessive gaps, and pauses break the accumulation. Gross movement is a total, not a rate. If a separate movement-rate outcome is derived, its denominator will be observed covered time rather than the entire nominal session. Otherwise, two runs with different data loss could produce misleadingly different motion totals.

Fixed-axis components derived from rotation vectors are not Euler yaw, pitch, and roll. Their labels will retain the actual calculation. Head motion also does not directly measure swivel-chair rotation, because the head can move relative to the torso and chair. A seated avatar drawn at the scene center in Trial Review is a visualization convention, not evidence that measured head translation was zero.

Instruction counts will use actual software playback starts. A requested cue, a rejected model answer, and a prefetched clip are not delivered instructions. Opening cues, search guidance, acknowledgments, reminders, and success speech will have separate categories. The planned total of delivered active-search assistance includes directional guidance, gaze/pointing acknowledgments, and selection reminders; each category will also be reported separately. Opening and success speech remain outside that total. Interrupted speech counts as started speech while retaining its completion status; it should not be treated as if the participant heard every word.

Sequence-Level Latency

For a qualified event, let t_q denote the earliest supported qualification time and t_p the actual software playback start of the corresponding acknowledgment. The primary engineering delay is

$$L_{\text{software}} = t_p - t_q. \tag{5.5}$$

This is not acoustic latency at the participant's ear. Audio-device buffering and the physical sound path require a separate measurement. Gaze and controller qualification times must also use their own definitions; comparing a controller acknowledgment against first gaze contact would inflate or misattribute its delay.

Repeated requests about one continuously qualified event form a single episode. A sequence can contain several Jev waits without creating several independent missed confirmations. An episode

ends with playback, disqualification, trial termination, or a defined reset. Evaluation will report how those endpoints were assigned and mark incomplete evidence as unknown.

Where logs support attribution, waiting will be divided into model response, explicit model deferral, scheduling, existing speech, and audio readiness. These intervals may overlap. They should not be added mechanically unless represented as a nonoverlapping partition of the critical path. The full event timeline remains the authority when aggregate attribution is ambiguous.

Figure 5.5 separates intentional evidence qualification from subsequent delivery delay. A model response can be timely while playback is delayed elsewhere, and overlapping waits must not be double-counted.

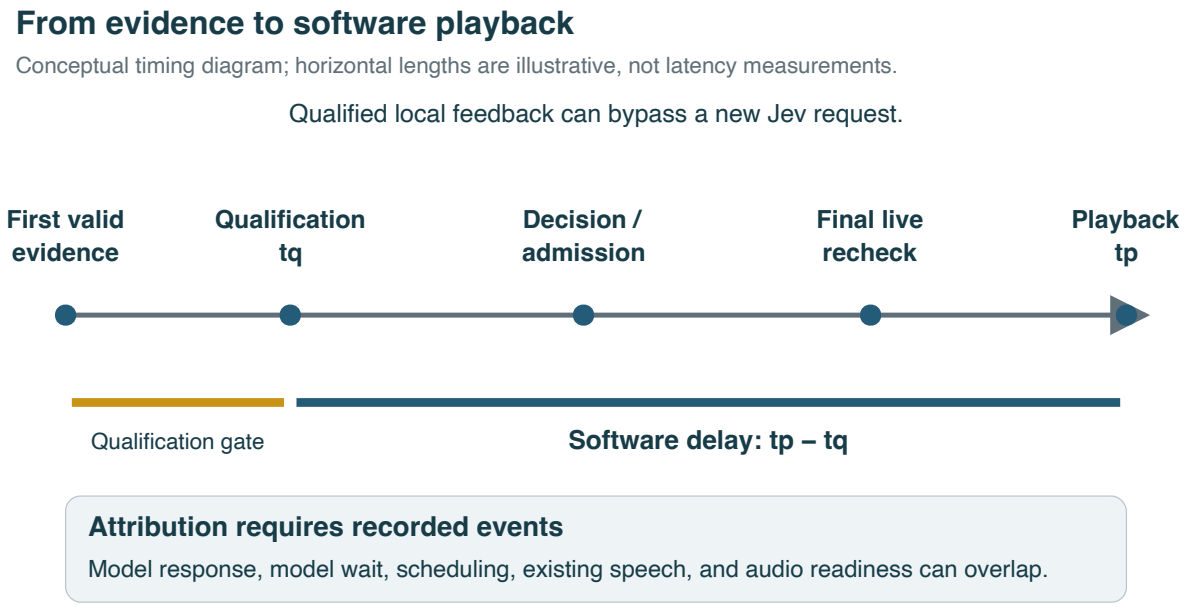


Figure 5.5: Conceptual event timeline, not a measured trace. Software delivery delay begins at supported qualification and ends at actual software playback start. The preceding qualification gate and the unmeasured acoustic path are separate quantities. Qualified local feedback can avoid a new cloud request.

Table 5.1 keeps the principal timing variables separate. In particular, the onset of an audio clip is a software event; neither a candidate phrase nor a successful provider response proves that the participant heard it.

Table 5.1: Timing contract for planned analysis and engineering evaluation. All cross-clock comparisons require an established mapping; unavailable evidence is not assigned a zero duration.

Quantity	Endpoints	Interpretation limit
Active completion	Reveal to correct trigger, less union of pauses	Excludes readiness and opening cue; unfinished trials remain unfinished.
First gaze contact	Reveal to first valid target-ray hit on active clock	Intersection, not fixation or vicinity acknowledgment.
First pointing contact	Reveal to first tracked right-ray hit on active clock	Not the one-second hold or trigger selection.
Jev response latency	Request dispatch to response receipt	A response can be deferred, canceled, or stale.
Confirmation delivery	Qualified evidence to matching actual software playback	Separate gaze and pointing episodes; not acoustic latency.
Live transfer lag	Source event to viewer receipt, if clocks are mapped	Describes transport, not headset decision delay.

Questionnaire Joins and Versioning

Final questionnaire administration is separate from the block to which a response refers. A response collected after both blocks will retain final-administration timing and a Block 1 or Block 2 reference. The actual voice identifier must be joined through the saved order, not inferred from block number. Practice is excluded from measured block outcomes even though it uses the first voice.

A convenience export may repeat a block-level workload score on every trial row. That does not create 40 independent workload observations. The analytical dataset will identify participant-by-block scope and collapse duplicate joins when estimating questionnaire effects. Raw responses,

score direction, scale bounds, item identifiers, instrument version, and missingness will remain available.

The current NASA-TLX adaptation uses 0–10 ratings, whereas older development exports used 0–100. Automatic scoring must require explicit range metadata. A separately named normalized score can support comparison, but raw values must remain unchanged. Unknown legacy performance anchors cannot be repaired by guessing their direction. Complete-subscale rules for IMI and agent-perception measures will likewise be implemented and tested before confirmatory analysis.

CHAPTER 6: PLANNED STATISTICAL ANALYSIS

Primary contrast

The primary estimand will compare active completion time under the self-similar and participant-preferred catalog voices, standardized to the same starting-difficulty distribution. A trial-level mixed-effects analysis of log completion time will account for participant clustering, voice condition, period/order, trial progression and absolute pre-cue starting theta. The analysis will report the voice time ratio and uncertainty using a two-sided contrast.

A prespecified low-complexity function of theta will allow nonlinearity; nominal-angle strata will support balance plots and a categorical sensitivity analysis. Zero-degree trials will remain in the analysis, and completion time will not be divided by angle. The model's random-effects structure, angle function and smallest meaningful effect will be finalized with simulation before confirmatory recruitment. Participant count will be justified under that same model; repeated trials will not be treated as independent participants.

Figure 6.1 illustrates the planned units of observation. Trial-level measurements are nested within participants, while retrospective ratings remain participant-by-block observations.

Secondary and exploratory outcomes

Retrospective condition-specific ratings will be collected only after both blocks, with actual survey timestamps retained. Reference scope remains participant-by-block. Recall and recency effects will be considered when interpreting condition differences; these are not immediate post-block measurements. Individual workload dimensions and complete unweighted workload means will

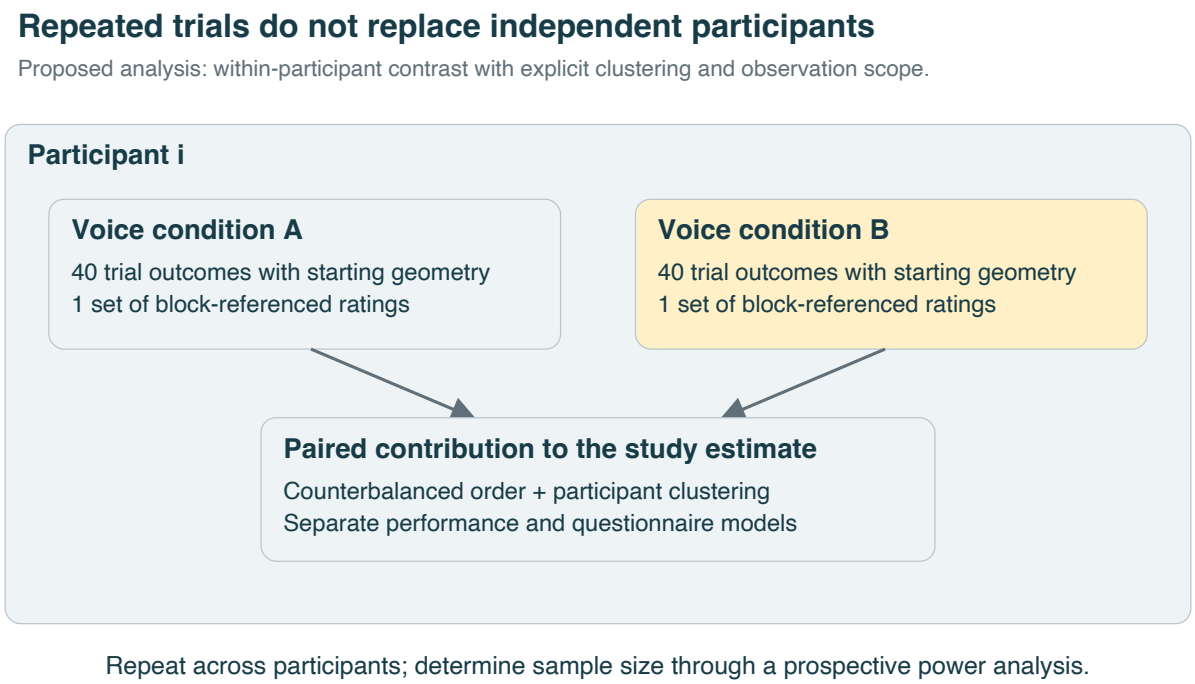


Figure 6.1: Proposed analysis scope. Each participant contributes repeated trial outcomes and two block-referenced questionnaire sets. Counterbalanced order and participant clustering remain part of the analysis plan; additional trials do not replace a justified participant sample size.

be summarized by participant and block. IMI subscales and the three six-item agent-perception constructs will remain separate scores. The additional Guo Q20–22 appealing, eerie and familiar ratings will be analyzed individually, without a combined eeriness score. Voice similarity will be a manipulation check. Custom experience ratings and self-reported guidance use will be analyzed individually. Multiplicity families and secondary contrasts will be specified before confirmatory analysis.

Instruction count, observed gaze path and other behavior potentially caused by voice condition will not automatically be covariates in the primary model. Voice-by-theta variation, selected catalog voice, acoustic quality and guidance timing will be exploratory or sensitivity analyses with their assumptions reported.

Table 6.1 summarizes the proposed outcome hierarchy. It is an analysis plan rather than a results table; final contrasts, missing-data handling, and multiplicity families require preregistration.

Table 6.1: Outcome hierarchy and observation scope. Engineering metrics diagnose the guidance system and do not replace participant outcomes.

Role	Unit	Planned interpretation
Primary performance	Trial within participant	Voice contrast in active time, adjusted for starting geometry and design factors.
Secondary experience	Participant by referenced block	Workload and separate IMI subscales, retaining paired scope.
Exploratory perception	Participant by referenced block	Agent constructs, individual custom items, and voice-similarity manipulation check.
Safety	Participant over full session	Baseline-to-final symptoms and stops; no separate voice-specific SSQ effect.
Engineering quality	Qualified event episode / request	Delivery delays, errors, repetition, and provider reliability with distinct denominators.
Qualitative experience	Participant response / coded excerpt	Debrief themes with an audit trail, not causal performance estimates.

Incomplete and interrupted observations

Raw attempts, incorrect selections, pause intervals, timeouts, technical failures and withdrawals will be retained. Stopped trials will not be coded as successful completion times. Timeout/censoring rules and the missing-data estimand will be specified before collection; slow trials will not be removed solely for being slow. Missing responses will not be replaced with zero. Questionnaire scores will follow the complete-subscale rule, with item-level missingness reported. Disliking or failing to recognize a cloned voice will not by itself exclude a participant.

Guidance and qualitative evaluation

Engineering analysis will distinguish event qualification, decision latency, model deferral, scheduling delay and actual playback. Gaze and controller confirmation will be evaluated separately. Sequence-level review will include what happened before and after a cue. Automatic checks and ratings produced by an offline language-model evaluator (Chapter 7) will remain advisory engineering evidence, not participant perception scores.

Debrief responses will be coded with a documented thematic process and an audit trail, with independent review of a subset. Perceived performance, voice comfort, reliance and possible own-voice use cases will be interpreted alongside, but not substituted for, measured task performance.

A Candidate Primary Model

A candidate model makes the intended estimand explicit while leaving the final specification to preregistration. For participant i, period b, and trial j, let $Y_{ibj} = \log T_{ibj}$ for successfully completed active-search intervals. A starting specification is

$$Y_{ibj} = \beta_0 + \beta_V V_{ib} + \beta_B B_b + \beta_O O_i \tag{6.1}$$

$$+ f(|\theta_{0,ibj}|) + g(j) + u_{0i} + u_{Vi} V_{ib} + \epsilon_{ibj}. \tag{6.2}$$

Here V indicates self-similar rather than preferred voice, B identifies the second period, and O identifies assigned voice order. The functions f and g represent prespecified low-complexity effects of exact starting angle and trial progression. In this candidate specification, j indexes position within a block (1–40); the period term represents the between-block change. The exact pre-cue angle requires the snapshot described in Chapter 5, which remains an implementation gap. This

model cannot be fitted as specified using nominal wall angles relabeled as exact measurements. Participant intercepts and, if supported by the design and model stability, participant voice slopes represent clustering and variation in within-participant contrasts. With only one block per voice, the slope variation can also absorb block-specific session effects; it is not a pure measure of stable individual voice responsiveness. The final random-effects structure is an open decision, informed by the experimental design and simulation rather than selected to obtain significance [4].

The exponentiated voice coefficient, $\exp(\beta_V)$, is a conditional multiplicative time contrast under this model. It will not automatically equal a marginal arithmetic mean ratio. The primary estimand remains the marginal contrast standardized over the prespecified common starting-difficulty distribution. The conditional coefficient will be reported separately as a model parameter, not substituted for that estimand. Predictions on the original seconds scale require an explicit retransformation procedure and uncertainty estimate. No fitted coefficient is available in this first pass.

A two-period design supports adjustment for period and sequence, but cannot identify every possible asymmetric carryover mechanism independently. For example, using a voice during practice could affect familiarity in its first block, and the second block could benefit from task learning. The report will describe this limitation rather than claim that adding an order term removes all carryover.

Sample Size and Precision Planning

Forty trials per block increases within-participant information, but does not turn one participant into 80 independent people. A power analysis must reflect the planned trial-level model, between-participant heterogeneity, within-participant correlation, geometry distribution, incomplete trials, and the effect size worth detecting. An old calculation for a different number of conditions cannot justify this design.

The proposed simulation workflow will generate participants with both voice blocks and the intended nominal-angle quotas. It will vary plausible residual variation, participant intercept variation, voice-slope variation, and missingness. The same model intended for the final analysis will be fitted to each simulated dataset. Convergence, bias, interval coverage, and power or precision will be summarized across candidate enrollment sizes. The chosen smallest meaningful time ratio and assumptions will be documented before recruitment. Synthetic pilot values used to test the workflow will be labeled as such and will not be presented as empirical variance estimates.

A retained pilot can inform variance assumptions only under a documented rule about whether those observations enter the confirmatory sample. Reusing pilot outcomes after extensive policy tuning risks treating development data as untouched confirmation. The study will preserve a build and protocol identifier so that a pilot collected under a different guidance policy can be separated from the frozen release.

Data Screening and Missingness

Data screening will first establish structural integrity: unique participant/session/block/trial keys, consistent voice assignments, chronological reveal and selection events, and required geometry provenance. It will then assess observation coverage and protocol deviations. This order avoids excluding a slow trial merely because it looks unusual while overlooking a malformed timestamp or an incorrect condition join.

Tracking loss can affect guidance and measurement simultaneously. A trial with sparse gaze coverage may have valid selection timing but unreliable gaze-contact latency. Exclusion should therefore be measure-specific where justified. A missing gaze contact does not automatically invalidate com-

pletion time, and a complete motion stream does not repair missing voice playback events. Each analysis will report its denominator and reasons for missingness.

Stopping, timeout, withdrawal, and technical failure require a prespecified treatment. A complete-case time model estimates completion times among observed successful trials under its assumptions. It may not describe all assigned trials when incomplete outcomes differ by voice. The final plan must define a complementary completion-probability or time-to-event analysis if incompletion is consequential. A timeout threshold has not been frozen in the source protocol and is not invented here.

Secondary Outcomes and Multiplicity

Workload dimensions, workload mean, and IMI subscales will be treated as specified outcomes at participant-by-block level. The analysis will retain paired observations and account for period where appropriate. The choice between ordinal-item and approximately continuous scale models will be made before confirmatory analysis and justified for each measure. Reporting only the items with favorable differences would distort the battery's intended scope.

Multiplicity families and adjustment rules remain to be finalized. Holm's procedure is available for a prespecified family of confirmatory secondary comparisons [18]; it is not a substitute for deciding which outcomes belong in that family. Exploratory agent-perception and custom-item results will be labeled, with effect estimates and uncertainty rather than an indiscriminate list of significance claims.

Voice similarity is a manipulation check. Failure to recognize the clone is informative about the manipulation and will not automatically remove a participant. Likewise, low liking is an outcome, not a reason to exclude an inconvenient response. Sensitivity analyses may investigate synthesis

quality or recognition, but conditioning on post-assignment experience changes the interpretation and must be stated.

Behavior as Outcome Rather Than Automatic Adjustment

Instruction count, gaze path, and controller behavior can be consequences of voice condition. Including them automatically as primary-model covariates could remove part of the effect the study intends to estimate. The primary comparison therefore adjusts for pre-cue difficulty and design variables, while guidance and behavior enter exploratory or sensitivity analyses under explicit causal assumptions.

For example, slower search can produce more opportunities for guidance, and more guidance can itself affect search time. A correlation between instruction count and duration does not establish that speech caused delay. Timeline analysis can identify concrete scheduling failures, but observational associations among adaptive behaviors require more cautious interpretation.

Qualitative Analysis and Mixed Evidence

The debrief will begin with broad experience questions, followed by balanced probes about voice preference, perceived performance, comfort, and potential uses. The coding process will record the initial codebook, revisions, coding decisions, and review of a subset. Contradictory or minority accounts will be retained rather than forced into one favorable theme.

Reported reliance and observed behavior will be considered together without treating either as a perfect measure of the other. A participant may say they relied on the agent while rapidly anticipating several cues, or may remember one frustrating interruption despite mostly timely guidance.

Such differences can guide interpretation but do not establish dishonesty or measurement failure. The final report will separate quotations, coder interpretation, and behavioral evidence.

Safety Outcomes and Manipulation Checks

Baseline and final Simulator Sickness Questionnaire responses will describe symptom burden and pre-to-post change across the full exposure. Canonical scores will be calculated only from the intended item anchors and scoring rules. These measurements do not provide separate voice-block sickness effects because there is no between-block administration. Safety stops, discomfort reports, and incomplete exposure will be reported alongside symptom change; a numerical exclusion or stopping cutoff is not invented for this draft. The operational stopping procedure must be frozen before participant collection.

The similarity manipulation check will report the paired difference between self-similar and catalog-voice ratings with uncertainty. A group-level decision rule and any practically meaningful difference remain open preregistration decisions. The study will not classify an individual as invalid merely because the clone is not recognized. Equivalence testing is not currently specified; a nonsignificant voice contrast will therefore not be interpreted as equivalence.

CHAPTER 7: ENGINEERING VERIFICATION AND RESULTS FRAMEWORK

Evidence Available for This Draft

This chapter distinguishes existing engineering records from participant results that have not yet been collected. The October 20 alignment audit documents software checks, analysis tests, and an Android build/install/launch. These records support specific implementation claims. They do not establish a treatment effect, a naturalness rating, eye-tracker accuracy, or acoustic response latency.

The audit reports passing schedule and voice-isolation checks, including one practice trial in the first assigned voice, two measured blocks, and nominal-angle quotas. It also records 38 trial-export tests and seven workload scoring/plot tests. Those counts describe the audit's test run, not a newly executed suite for this manuscript. The verification record should accompany any later thesis statement so the tested revision remains identifiable.

Why Sequence Evaluation Is Necessary

A phrase can be geometrically admissible in isolation and still be poor guidance in sequence. Repeating a direction immediately after it was delivered can be unnecessary. A correct downward cue can become obsolete after the controller reaches the target. A gaze acknowledgment can arrive late after multiple model waits even though each individual response was well formed. Evaluation must therefore inspect recent state, actions, and actual speech around the event.

The deterministic checks will establish whether the evidence supports an action. They will distinguish gaze confirmation, controller confirmation, and trial completion. A model judge may then

Table 7.1: Evidence status at the documented implementation checkpoint.

Evidence	Supported conclusion	Boundary
Schedule and voice checks	One practice and nominal 40-by-two allocation pass fixtures	Does not verify wearer completion of a full session
Trial export and TLX tests	Declared scales and retrospective joins pass tested cases	Not validation of all questionnaires
Unity integration checks	Listed state and interaction fixtures pass	Not proof against all timing races
Android build/install/launch	Application built, installed, and started	Not a participant acceptance session
Exact pre-cue theta	Required definition documented	Snapshot implementation remains a gap
Vertical allocation	Cross-relative definition documented	20/20 enforcement remains a gap
Voice-condition outcomes	No confirmatory dataset reported	No performance or preference conclusion

assess wording, timing, repetition, and interruption appropriateness using that evidence. Its ratings remain advisory. They are not human-perceived correctness and do not convert an incompletely logged episode into a known success or failure.

Paired Policy Pilots

Paired pilots will feed identical recorded states and scripted trajectories to the baseline and candidate policies. The scenario set includes large turns, wall entry, vertical refinement, overshoot, gaze on target with the controller elsewhere, controller on target with nearby gaze, overlapping vicinity regions, blinks, stale tracking, pauses, and delayed model responses. Audio occupancy and readiness must also be represented so that playback admission can be compared fairly.

The candidate-policy comparison may include small, medium, and full eligible phrase sets. Full eligible means all choices that still satisfy hard geometry and validity requirements; it does not

mean handing the model an unfiltered catalog and ignoring physical constraints. Provider support and exact candidate counts will be recorded. Experimental policies will remain outside participant builds until their results justify a controlled change.

Replaying an observed trajectory is a counterfactual engineering test. The original participant did not hear the alternative sequence, so the replay cannot demonstrate that the alternative would have reduced their search time. Scripted trajectories provide reproducible edge cases but are not measured human behavior. Both are useful for identifying faults before a wearer pilot.

Metrics and Episode Definitions

For each policy, the report will summarize qualification-to-software-playback latency with median, 95th percentile, maximum, and count beyond a stated engineering threshold. The threshold must be declared before comparison and distinguished from a perceptual or clinical criterion. Episodes, rather than repeated requests within an episode, will be the denominator for delayed confirmations.

Jev request latency, failures, timeouts, calls per trial, cancellations, and late responses will be reported separately. A provider timeout can lead to a timely local fallback, while a successful provider response can still arrive too late for the current state. Provider reliability and delivered guidance reliability therefore require different measures.

Wrong-direction and stale hints will be labeled only when the available evidence shows the instruction is unhelpful toward the task. A different but valid phrase is not automatically wrong. Missing evidence remains unknown. Repetition, unnecessary speech, excessive silence, and interruption appropriateness require context beyond a sign check on an angle.

Illustrative Failure Analysis

An ordering failure can occur when a correction is queued before the controller reaches the target. If playback checks use the older snapshot, the correction can play immediately before a pointing acknowledgment. Another possibility is that the controller hold has not qualified yet and the ordinary-direction policy remains active during the hold. A third is that action priority admits a lower-priority cue despite newer pointing evidence. These causes require different fixes and cannot be distinguished from the spoken text alone.

The intended correction is to suppress contradictory away-directions as soon as valid controller-on-target evidence is present, while retaining the independent qualification timer for the spoken pointing acknowledgment. A final playback check must repeat the suppression after any asynchronous wait or audio fade. If the controller leaves the target before qualification, no pointing acknowledgment should be emitted merely because the earlier request was valid.

A separate reference-frame problem arises when an object appears higher than a highlighted neighbor but the instruction says to look lower. The instruction may be calculated relative to the current eye ray, which could lie above both. The phrase can then be geometrically consistent yet confusing because the highlight suggests a different reference object. Diagnosis requires the recorded gaze ray, object positions, highlighted candidate, and playback time. It must not rely only on the review camera's screen-space layout.

These examples are diagnostic categories drawn from the development history. This chapter does not report their incidence or claim they have all been eliminated in a frozen participant build. The acceptance matrix in Appendix C defines evidence required to make those claims.

Participant Results to Be Added

The completed empirical chapter will first report participant flow, exclusions, incomplete sessions, and data coverage by condition and order. It will then describe realized starting geometry and the similarity manipulation check before presenting the primary voice contrast. This order allows readers to assess whether the intended comparison was delivered and measured.

The primary result will include the fitted time contrast, its uncertainty, the analyzed participant and trial counts, and model diagnostics. Secondary sections will report workload and IMI subscales at participant-by-block level. Exploratory sections will report agent perception, voice experience, guidance behavior, and debrief themes. Tables will identify missing observations rather than silently change denominators.

Table 7.2: Planned empirical results, not observed findings.

Result family	Required evidence	Present status
Participant flow	Consented, started, completed, stopped; reasons	Not available
Manipulation check	Similarity by voice with paired scope	Not available
Primary performance	Adjusted time contrast and uncertainty	Not estimated
Workload and IMI	Versioned raw items and complete subscales	Not collected for this report
Agent experience	Individual and construct-specific ratings	Not collected for this report
Guidance quality	Sequence-level playback and evidence	Evaluation design specified
Debrief themes	Recorded responses and coding audit	Not available

Interpreting Two-Second Delays

A two-second delay can contain several different intervals: waiting for qualification, waiting for a model response, an explicit model decision to wait, queued speech, audio readiness, and playback scheduling. Removing a provider wait does not establish that all delays disappear. Conversely, a one-second controller hold is an intentional qualification period, not necessarily a one-second processing fault.

The appropriate test starts from the first supported qualification for the particular event and ends at its actual software playback. A gaze event with a 100-ms gate and a controller event with a one-second hold must not share an undifferentiated origin. The report will retain both first evidence and qualification time so that readers can distinguish intentional gating from added delivery delay. Without a new verified sequence dataset and a wearer or acoustic measurement where required, this draft cannot claim that the roughly two-second delays have been removed.

CHAPTER 8: DISCUSSION AND LIMITATIONS

What the Study Can Explain

The proposed study can compare two voice experiences within the same adaptive search system. Because the catalog voice is chosen by the participant, the comparison asks whether self-similar synthesis differs from a preferred alternative. It does not ask whether personalization is better than an arbitrary default voice. This choice strengthens the practical relevance of the comparison while narrowing the causal claim.

The system also makes the assistance process observable. A slower trial can be examined alongside starting geometry, tracking availability, controller behavior, and delivered speech. That evidence can distinguish a long search with timely guidance from one affected by stale cues or a silent interval. It cannot by itself determine why a participant trusted or ignored an instruction. The subjective and qualitative measures address aspects of experience that the event stream does not directly observe.

Possible Outcome Patterns

If self-similar voice produces faster active completion with acceptable uncertainty, the interpretation will still depend on the manipulation check and guidance quality. A difference accompanied by systematic intelligibility problems in the catalog condition would not isolate resemblance cleanly. If both conditions have comparable delivery quality and participants recognize the clone, the result would more directly support a voice-condition effect in this task.

If the preferred catalog voice produces faster completion, that would be consistent with the possibility that a selected external voice is easier to follow, but the study would not establish the mech-

anism without further evidence. Liking, naturalness, and reliance ratings may help describe the experience. They should not be treated as proven mediators solely because they correlate with performance.

If completion times are similar, uncertainty matters. A wide interval would leave practically important differences unresolved. A narrow interval near the null could constrain the size of a plausible effect under the tested conditions. A nonsignificant test alone will not establish equivalence. Equivalence testing is not part of the current analysis specification. Adding it would require a meaningful bound and procedure defined before inspecting the result.

Experience and performance may diverge. Participants may prefer their cloned voice while searching no faster, or may find it eerie while completing efficiently. Such outcomes would be substantively useful because they distinguish usability, affect, and speed. The battery is designed to retain these differences rather than force a single ranking.

Voice Identity and Acoustic Confounds

Voice condition includes more than a nominal label. The self-similar voice depends on recording quality and synthesis, while the preferred voice comes from a heterogeneous catalog. Both may differ in pitch, prosody, speaking rate, pronunciation, and clip duration. The same words and provider do not guarantee matched acoustics. The study will document these properties and pilot comfortable intelligibility rather than claim exact equivalence.

Participant choice is intentional heterogeneity. Requiring every participant to use the same catalog voice would answer a different question. The analysis should retain chosen voice identifiers and describe the distribution, while avoiding an overfitted model with unsupported estimates for sparsely

chosen voices. An exploratory association between a catalog voice and performance may reflect who selected it as well as its acoustic properties.

No extra pitch, panning, or inner-thought treatment is part of the current self-similar manipulation. Earlier attempts to make cloned speech sound like an internal monologue are development history. Reintroducing such an effect in only one condition would change the manipulation and require a revised rationale and analysis.

Adaptive Guidance as a Shared Mechanism

A shared policy does not produce identical spoken sequences. Participants turn, pause, look, and point differently. The policy consequently supplies different directions and can emit different numbers of hints. This behavior dependence is part of the design, not automatically a failure of experimental control.

Nevertheless, policy drift would threaten the comparison. Changing phrases, timers, or provider behavior midway through collection could alter assistance quality. The study release must pin the application build, configuration, phrase semantics, voice preparation settings, and decision-provider version where possible. Any unavoidable service change must be documented and considered in the analysis.

The central state machine improves auditability by giving one local owner to guidance evidence and transitions. It does not guarantee perfect guidance. Calibration error can produce an incorrect observation; an unclear reference frame can produce an unhelpful phrase; a stale response can escape an incomplete check. A defensible report describes the safeguards and measured failures rather than implying that architectural separation proves correctness.

Practice, Learning, and Order

One practice trial reduces setup burden but provides limited exposure to a complex interaction. A participant must learn to hold gaze on the readiness cross, follow speech, use the chair comfortably, aim the right controller, and trigger selection. The practice-proficiency rule remains to be established through pilots. A fixed single trial does not prove that every participant has acquired these skills equally.

Practice uses the first assigned voice. This aligns initial exposure with the first block, but also couples practice familiarity with period. Counterbalancing helps distribute this across voice conditions. It does not prevent fatigue, learning, novelty, or asymmetric carryover. These limitations will be reported with period and order effects and any relevant sensitivity analysis.

Retrospective Questionnaires and Burden

The final-only survey avoids interrupting the two-block search sequence with a long questionnaire. It also asks participants to recall Block 1 after completing Block 2. Ratings may therefore reflect recency, contrast, or imperfect recall. The two sets will explicitly identify the referenced block and voice, but that labeling does not guarantee accurate recollection.

The condition-specific battery contains 130 numeric responses across both blocks, in addition to symptoms, setup ratings, comparison questions, and debrief. Completion time, missingness, and comprehension require pilot measurement. A participant who is tired or symptomatic should first receive the immediate post-exposure safety check. The survey's breadth is a current design decision, not evidence that every item is necessary or equally interpretable.

Adapting an embodied-agent measure to a voice agent can alter its meaning. Co-presence or anthropomorphism items may be understood differently without a visible character. The thesis will retain attribution and wording changes, report the adaptation, and avoid asserting new psychometric validation from a small experimental sample. Custom reliance and naturalness items remain individual exploratory measures.

Gaze Quality and Feedback Coupling

Gaze serves both as an input to assistance and as a recorded behavioral signal. Poor tracking can therefore change the participant's experience while also reducing measurement coverage. This coupling complicates interpretation: lower gaze availability is not merely a missing column if it also suspends guidance. Coverage and technical events must be considered alongside performance without automatically treating them as pre-existing participant traits.

A highlighted vicinity candidate can help indicate what the system currently attends to, but it can also suggest a reference object that differs from the mathematical gaze ray. The known lower-versus-higher confusion illustrates this problem. A cue should be evaluated in its visible interaction context, not solely against the sign of a geometric delta. Development-only rings and diagnostic markers must not silently enter the participant condition.

The proposed validity-aware stream and independent calibration checks address these risks, but headset-specific spatial performance and distinct acquisition rate remain empirical questions. The report must not borrow a Vive Pro Eye validation result and label it Focus Vision performance. Nor should it infer cognition from a convenient interaction threshold.

Privacy, Consent, and Reproducibility

A coded participant identifier does not make a voice recording anonymous. Natural and synthesized speech can remain recognizable. The consent and data-management documents must describe recording, synthesis, cloud processing, storage, access, and deletion using the services actually employed. Older local-only descriptions are inconsistent with current ElevenLabs voice preparation and Jev decisions and require reconciliation before participant use.

The thesis will report the relevant processing stages without including credentials or identifiable recordings. A reproducibility package can provide code, versioned schemas, example synthetic records, and de-identified summaries where authorized. Distribution of cloned voices or raw participant audio requires a separate basis and must not be assumed from consent to participate.

Cloud dependencies also affect replication. A future researcher may not obtain the same voice model or identical decision behavior. Preserving prepared-library metadata and actual playback records helps describe the original exposure, while a constrained phrase interface makes alternative providers testable. It does not eliminate the need to report provider changes.

Future Work

A future study could isolate memory, guidance policy, or degree of model autonomy as independent factors. Such a study would require a new experimental design rather than reinterpret the current voice contrast. Provider comparisons could begin with paired replay and then use wearer pilots to test whether observed scheduling improvements translate to a better interaction.

Other extensions include a visible agent body, a known-target search condition, free-walking navigation, a larger range of voice choices, and accessibility-oriented selection alternatives. Each

changes the task or population. The present seated right-controller design supplies a controlled starting point, not a universal model of mixed-reality assistance.

Finally, validated eye-event and pupil analyses could enrich the behavioral account if acquisition, calibration, lighting, and event classification support them. They should be added because they answer a defined question with valid measurements, not because the headset exposes a convenient data field.

CHAPTER 9: CONCLUSION

Follow My Voice connects a personalized voice comparison to an instrumented mixed-reality search task. The participant does not know the target identity in advance. Instead, gaze-contingent speech guides orientation and within-wall refinement, while the right controller provides pointing and explicit trigger selection. This organization makes the agent relevant to the search without treating gaze as a selection command.

The system separates spatial evidence, state transitions, constrained model choice, scheduling, and delivered speech. That separation is important when evaluating delayed or contradictory guidance: a model proposal, a qualified gaze event, a controller hold, and a successful selection are different observations. Recording them distinctly makes the interaction more auditable and supports a more precise account of what went wrong when assistance is unhelpful.

The proposed study compares a preferred catalog voice and a self-similar synthetic voice in two counterbalanced 40-trial blocks, following one practice trial in the first assigned voice. Its primary outcome is active completion time adjusted for starting geometry and participant clustering. Retrospective workload, motivation, and agent-experience measures will complement performance after both blocks. The comparison will concern voice identity within a shared adaptive policy, not an isolated benefit of gaze contingency or memory.

At this stage, the contribution is the system, operational design, and measurement plan. Confirmatory participant results are not available, and the thesis makes no claim that either voice improves performance. Before collection, the study must resolve exact pre-cue geometry, balanced vertical allocation, sample-size planning, pilot acceptance criteria, and consent alignment with cloud processing. The completed thesis will then be able to connect its conclusions to actual participant evidence, with uncertainty and limitations stated alongside any observed differences.

APPENDIX A: QUESTIONNAIRE WORDING AND ADMINISTRATION

This appendix preserves the current Overleaf instrument wording. Administration remains proposed, with Qualtrics synchronization on hold. The two retrospective block references do not imply questionnaires between blocks. Source attribution and scoring rules are described in the methodology chapter.

Administration

Two counterbalanced voice blocks will contain 40 measured trials each, with one separate practice trial. The final survey will administer two separately labelled retrospective questionnaire sets, one referring to each voice block, using identical item wording and order. No questionnaires will be administered between blocks. The break and immediate comfort/tracking checks will remain. Surveys and rest breaks will occur outside active-search timing.

Stage	Measures
Before headset exposure	Demographics, experience and baseline SSQ
Voice setup	Agent-gender preference, audition ratings and voice choice
Between voice blocks	Rest break, immediate comfort and tracking checks; no questionnaire
Final survey after both blocks	Post-exposure SSQ; two condition-specific sets (6 workload, 24 IMI, 21 agent-perception and 14 voice/assistance ratings each), technical checks, optional comments, voice comparison and debrief

Each condition-specific set in the final survey contains 65 numeric responses (130 across both sets) before technical checks and optional text. Setup will include a separate liking rating for each auditioned voice, including the self-similar voice. Baseline and final SSQ together contain 32 symptom ratings.

Baseline

`BASE_STUDY_CODE`

Study code (assigned by the researcher): *Response:* Coded identifier; no name or contact details.

`BASE_AGE`

Age range: *Response:* 18–24; 25–34; 35–44; 45–54; 55–64; 65+; Prefer not to answer.

`BASE_GENDER`

Gender identity (select all that apply): *Response:* Woman; Man; Nonbinary; Another identity (optional description); Prefer not to answer.

`BASE_HANDEDNESS`

Which hand do you usually use for everyday tasks? *Response:* Right; Left; Both equally; Prefer not to answer.

`BASE_CORRECTION`

Do you use glasses or contact lenses? *Response:* Neither; Glasses; Contact lenses; Both; Prefer not to answer.

`BASE_MR`

How often have you used virtual- or mixed-reality headsets? *Response:* Never; Once or twice; Several times; Monthly; Weekly or more.

`BASE_EYE`

Have you used eye-tracking technology before? *Response:* No; Yes; Unsure.

`BASE_VOICE`

Have you used a synthetic or cloned version of your own voice before? *Response:* No; Yes; Unsure.

`BASE_AI_FAMILIARITY`

How familiar are you with AI voice assistants? *Response:* 1 = Not at all familiar; 2–6; 7 = Extremely familiar.

Current symptoms

`BASE_SSQ`

Rate each symptom now, before headset use. *Response:* One severity rating per symptom.

1. General discomfort
2. Fatigue
3. Headache
4. Eyestrain
5. Difficulty focusing
6. Increased salivation
7. Sweating
8. Nausea
9. Difficulty concentrating

10. Fullness of head
11. Blurred vision
12. Dizziness with eyes open
13. Dizziness with eyes closed
14. Vertigo
15. Stomach awareness
16. Burping

Response: 0 = None; 1 = Slight; 2 = Moderate; 3 = Severe.

Voice setup

`SETUP_AGENT_GENDER`

What gender presentation would you prefer for the agent's voice? *Response:* Masculine; Feminine; Another presentation (optional description); No preference; Prefer not to answer.

`SETUP_CANDIDATE_LIKING`

How much do you like this voice? *Response:* 0–100, with 0 = Do not like at all and 100 = Like extremely; Unable to rate. Repeat for every auditioned candidate and the self-similar voice.

`SETUP_CHOSEN_VOICE`

Which of these voices would you prefer to use for the task? *Response:* Select from the voices actually offered. Ties in liking ratings will be resolved by this explicit choice.

`SETUP_AUDIO_CHECK`

Can you hear the voice clearly and at a comfortable volume? *Response:* Yes; No. Resolve volume or audibility problems before continuing. Record the offered catalog/version, voice IDs, audition order, selected voice and any unavailable preference. Demographic gender, preferred agent presentation and perceived voice gender are distinct variables. No preference will be inferred from demographic gender.

Final survey: condition-specific questionnaire sets

Complete this set twice during the final survey, first for Block 1 and then for Block 2. Display the referenced block number and actual voice condition before each set. Please answer only about that indicated block. There are no right or wrong answers. You may leave an item unanswered if you cannot answer it. The researcher will associate the response with participant, session and block IDs, actual voice condition/ID, condition order, survey version and actual administration timestamp. Record administration phase as final and reference scope as participant-by-block; the BLOCK item IDs describe the referent, not collection timing.

Workload

Rate each aspect of the indicated block.

1. Mental Demand – How mentally demanding was this block?
2. Physical Demand – How physically demanding was this block?
3. Temporal Demand – How hurried or rushed did you feel during this block?
4. Performance – How unsuccessful were you in accomplishing the task?
5. Effort – How hard did you have to work to achieve your level of performance?
6. Frustration – How insecure, discouraged, irritated, stressed, or annoyed did you feel?

Response: Integer ratings from 0 to 10. For demand, effort and frustration: 0 = Very low; 10 = Very high. For performance: 0 = Perfect; 10 = Failure. Intermediate values are numbered without additional anchors.

Activity experience: IMI

Rate each statement about the search activity in the indicated block. Use 1 = Not at all true, 4 = Somewhat true, and 7 = Very true; the other numbered points are available.

Interest/enjoyment

1. I enjoyed doing this activity very much.
2. This activity was fun to do.
3. I thought this was a boring activity.
4. This activity did not hold my attention at all.
5. I would describe this activity as very interesting.
6. I thought this activity was quite enjoyable.
7. While I was doing this activity, I was thinking about how much I enjoyed it.

Effort/importance

1. I put a lot of effort into this.
2. I didn't try very hard to do well at this activity.
3. I tried very hard on this activity.
4. It was important to me to do well at this task.
5. I didn't put much energy into this.

Pressure/tension

1. I did not feel nervous at all while doing this.
2. I felt very tense while doing this activity.
3. I was very relaxed in doing this.
4. I was anxious while working on this task.
5. I felt pressured while doing this activity.

Value/usefulness

1. I believe this activity could be of some value to me.
2. I think that doing this activity is useful for practicing visual search.
3. I think this is important to do because it can improve my search performance.
4. I would be willing to do this again because it has some value to me.
5. I think doing this activity could help me to search for targets.
6. I believe doing this activity could be beneficial to me.
7. I think this is an important activity.

Agent perception

Rate the voice agent you experienced in the indicated block. Use 1 = Very strongly disagree and 7 = Very strongly agree; intermediate points are numbered 2–6.

Co-presence

Q1. I noticed the voice agent.

Q2. The voice agent noticed me.

Q3. The voice agent's presence was obvious to me.

Q4. My presence was obvious to the voice agent.

Q5. The voice agent caught my attention.

Q6. I caught the voice agent's attention.

Perceived intelligence

Q13. The voice agent was able to operate without my intervention.

Q14. The voice agent was aware of the virtual environment.

Q15. The voice agent was able to set and pursue tasks by itself in anticipation of future needs.

Q16. The voice agent was able to complete tasks quickly.

Q17. The voice agent was able to find and process the necessary information for completing the task.

Q18. The voice agent was able to adapt/adjust its behavior based on prior events.

Appealing, eerie and familiar impressions

Q20. I found the voice agent appealing.

Q21. I found the voice agent eerie.

Q22. I found the voice agent familiar.

Perceived anthropomorphism

Q37. The voice agent is able to speak like a human.

Q38. The voice agent can be happy.

Q39. The voice agent can be friendly.

Q40. The voice agent can be respectful.

Q41. The voice agent can be funny.

Q42. The voice agent can be caring.

Voice and assistance experience

Rate each statement about this block. Use 1 = Strongly disagree and 7 = Strongly agree; intermediate points are numbered 2–6.

1. The agent's voice is similar to mine.
2. The hints seemed responsive to what I was looking at.
3. The hints were helpful.
4. The hints were distracting.
5. The speaking voice felt familiar.
6. I trusted the spoken guidance.
7. I felt comfortable hearing this voice.

8. The voice felt uncanny or unsettling.
9. I relied on the spoken hints when deciding where to search next.
10. The agent's voice sounded natural.
11. I was willing to follow the agent's instructions.
12. I was motivated to complete the task without the agent's guidance.

`BLOCK_VOICE_LIKING`

How much did you like the voice in this block? *Response:* 0–100; 0 = Do not like at all, 100 = Like extremely; Unable to rate.

`BLOCK_GUIDANCE_USE_PCT`

During this block, what percentage of the time did you use the agent's guidance to decide where to search? *Response:* 0–100%; 0 = None of the time, 100 = All of the time; Unable to estimate.

`BLOCK_PERCEIVED_GENDER`

How would you describe the gender presentation of the agent's voice? *Response:* Optional free text; Unsure.

Technical and comfort checks

`BLOCK_HEARD`

Did you hear all prompts clearly? *Response:* Yes; No; Unsure.

`BLOCK_FAULT`

Did any prompt unexpectedly repeat, cut off, arrive late, or use the wrong voice? *Response:* No; Yes; Unsure. If Yes, optional description.

`BLOCK_BREAK`

Did you take an unplanned break or experience an interruption? *Response:* No; Yes. If Yes, optional description.

`BLOCK_COMMENT`

Is there anything else you would like to share about this block? *Response:* Optional free text.

Final survey: symptoms, comparison and debrief

Post-exposure symptoms and comfort

Administer these current-symptom items immediately after headset removal, before the retrospective condition-specific sets. Rate each symptom now.

1. General discomfort
2. Fatigue
3. Headache
4. Eyestrain
5. Difficulty focusing
6. Increased salivation
7. Sweating

8. Nausea
9. Difficulty concentrating
10. Fullness of head
11. Blurred vision
12. Dizziness with eyes open
13. Dizziness with eyes closed
14. Vertigo
15. Stomach awareness
16. Burping

Response: 0 = None; 1 = Slight; 2 = Moderate; 3 = Severe.

`POST_STOP`

Did you stop or pause because of symptoms? *Response:* No; Yes.

`POST_READY`

Do you feel ready to leave the lab? *Response:* Yes; No. A No response requires researcher follow-up.

Voice comparison and debrief

`POST_VOICE_PREF`

Which voice did you prefer? *Response:* The catalog voice I selected; The synthetic version of my own voice; No preference; Unsure.

1. Why did you choose that response?
2. How did you approach the task, and how did the guidance affect your strategy?
3. What differences, if any, did you notice between the two voices?
4. Which moments were helpful, confusing, delayed or distracting?
5. Did either voice seem to affect your performance? What makes you say that?
6. What was it like hearing the synthetic version of your own voice? What did you like or dislike about it?
7. Did anything feel comfortable, unfamiliar or eerie? What contributed to that feeling, if anything?
8. In what situations, if any, would you prefer a voice like your own? When would you prefer another voice or no voice?
9. Did either voice feel like your own thoughts, like another person helping, or neither?
10. Do you have any concerns about how your synthetic voice was created or could be used?
11. What would you change about the guidance, voices or task?

Responses are optional. Ask the broad experience questions before the specific eeriness and inner-dialogue probes. Use follow-ups when needed rather than repeating a question already answered. Record perceived performance separately from logged performance.

Retained Voice-Recording Passage

The English passage below is retained verbatim from the current run sheet. Its match to the release build, recording duration, and audio acceptance procedure must be verified and versioned before participant use. Its object names are recording material and do not reveal the targets of measured trials. Participants will read in their usual speaking voice without a requested whisper or inner-thought performance.

> Locate the red sphere. Look toward the upper shelf. Move your gaze slowly to the right. Check the blue cube near the center. Compare each object's color and shape. Scan the lower row from left to right. Look beside the purple cylinder. Point the controller ray at the target and check its highlight. Press the trigger to select the matching object. Take a short pause. Get ready for the next trial.

APPENDIX B: DATA CONTRACT AND SCORING RULES

Record Identity and Provenance

The following dictionary specifies semantic requirements. Names shown here are descriptive contract labels and must be mapped to the versioned exporter schema before collection. They do not assert that every field already exists in the current build.

Table B.1: Required record identity and scope.

Field family	Meaning
Participant ID	Coded study identifier; not a name or voice-provider credential.
Session and run	Unique attempt/session context, including development versus study mode.
Block and trial	Referenced measured block, within-block trial, global trial index, and separate practice flag.
Voice condition	Preferred catalog or self-similar; actual voice ID, provider, synthesis model and configuration.
Assigned order	Preferred-first or self-similar-first, with practice following first assigned voice.
Build and policy	Application revision, configuration, state-machine policy, phrase semantics, and library version.
Schedule	Seed, schedule version, nominal wall assignment, realized geometry, and attempt status.
Survey scope	Administration phase and timestamp distinct from referenced participant/block/voice.

Geometry and Timing Fields

Table B.2: Geometry and event contract.

Quantity	Unit/reference	Required interpretation
Nominal angle	Degrees; wall schedule	Allocation category; not exact target bearing.
Exact starting theta	Degrees; pre-cue cross	Signed horizontal angle from saved origin-to-cross toward origin-to-target; currently a readiness gap.
Height difference	Metres; cross center	Target Y minus cross Y; intended and realized sign stored separately.
Target distance	Metres; first tracked head after reveal	Three-dimensional distance with observation timestamp.
Reveal	Session clock	Object visibility and active-search onset event.
First gaze contact	Active clock	First valid target collider intersection; missing if absent.
Last gaze contact	Active clock	Last valid observed target intersection; not fixation duration.
First controller contact	Active clock	First tracked right-ray intersection with target.
Gaze qualification	Session clock	Valid winning vicinity evidence meets interface timer.

Quantity	Unit/reference	Required interpretation
Pointing qualification	Session clock	Controller hold and current gaze requirements met.
Correct selection	Session and active clocks	Trigger selection of target; separate from acknowledgment.
Pause intervals	Session clock	Start, end, reason; union used for active time.
Instruction onset	Session clock	Actual software playback start, phrase and event category.

Raw Observation Requirements

Gaze records should retain available source time or sequence, application time, validity, source identity, origin, direction, and sample age. Per-eye data must be labeled as such when available; combined gaze must not be described as independent left and right observations. A head-direction fallback must not silently replace eye gaze in an eye-derived outcome.

Motion records should retain head and right-controller poses, tracking state, trigger state, and the time basis needed for reconstruction. Object records should retain stable identifiers, shape/color metadata, transforms, target designation, wall assignment, and layout version. Event records should identify state transitions, eligible actions, model requests and responses, cancellations, playback, interruptions, and trial outcomes. Unknown fields remain unavailable rather than receiving plausible-looking defaults.

Scoring Contract

NASA-TLX will retain six raw values on the documented 0–10 scale. Performance runs from Perfect to Failure, so larger values have the same adverse direction as greater demand, effort, and frustration. The unweighted mean requires all six. A normalized 0–100 field, if exported, must be named separately and calculated from explicit bounds. Legacy data without known bounds or performance direction remain unscored.

IMI scoring uses separate complete subscale means: interest/enjoyment (7), effort/importance (5), pressure/tension (5), and value/usefulness (7). On the 1–7 response scale, reverse-coded values use $8 - r$. Within-subscale reverse positions are enjoyment 3 and 4, effort 2 and 5, pressure 1 and 3, and none for value. There is no overall IMI total.

The selected agent-perception battery retains source identifiers. Co-presence Q1–6, perceived intelligence Q13–18, and anthropomorphism Q37–42 are separate six-item means when complete. Q20 appealing, Q21 eerie, and Q22 familiar are individual ratings in their original directions. There is no 21-item or combined eeriness total. Custom voice/assistance items likewise remain individual outcomes.

Export Validation Cases

Before analysis, automated checks should reject duplicate trial keys, unknown voice/block joins, negative active intervals, selection before reveal, and scores outside declared bounds. They should preserve incomplete outcomes, missing contacts, absent questionnaire items, and old records whose provenance is insufficient. A joined table must verify one survey observation per participant and

referenced block before repeating values on trial rows. Exports should be reproducible from raw records and a versioned analysis configuration.

APPENDIX C: GUIDANCE SEQUENCE TESTS AND EVALUATION RUBRIC

Controlled Scenario Matrix

Each scenario must save input observations, policy/configuration, eligible actions, provider behavior, playback occupancy, and actual admitted speech. Tests use the same trajectories across policies. The expected behavior below is an acceptance specification, not a claim that a wearer experienced the scenario successfully.

Table C.1: Sequence scenarios and acceptance criteria.

Scenario	Evidence and expected behavior
Large turn	Target opposite starting reference; progress cues remain directionally valid as wall boundaries are crossed.
Wall entry	Stable entry changes to refinement; looking upward on the same wall does not create a second arrival.
Vertical refinement	Target above/below gaze reference; magnitude wording follows documented plane-relative thresholds.
Overshoot and return	Valid recent crossing supports a return cue; a merely old target bearing does not.
Gaze only	Winning vicinity qualifies with controller elsewhere; gaze acknowledgment does not say the controller is on target.
Controller entry	Actual target intersection suppresses contradictory directions before hold qualification.
Controller hold	Qualified hold plus current valid nearby gaze admits pointing acknowledgment; no trigger completion is inferred.
Controller departure	Leaving before qualification cancels pointing acknowledgment and allows appropriate guidance after current-state checks.

Scenario	Evidence and expected behavior
Reacquisition	Valid gaze away rearms confirmation; continuously holding gaze does not repeat it indefinitely.
Overlapping regions	Only the closest eligible vicinity candidate receives acknowledgment; neighbor overlap alone is insufficient.
Blink or stale gaze	Missing evidence is not counted as a gaze hold or inferred departure; new gaze-dependent cues suspend.
Pause and recovery	Pending cues cannot leak across reset/pause; valid recovery resumes from fresh evidence.
Late Jev response	Previously valid direction is discarded if current controller/gaze state conflicts.
Provider wait/failure	Deferral remains bounded; fallback respects current eligibility and repetition limits.
Speech occupied	Higher-priority event follows explicit interruption/admission policy; interrupted text is not recorded as fully heard.
Correct trigger	Trial ends only on correct selection; success feedback uses completed-trial context.

Deterministic Labels

A wrong-direction label requires supported geometry showing that the delivered cue moves the user away from the intended target relation under the cue's reference frame. A stale label requires evidence that the relevant state changed before playback. Contradictory ordering requires a deliv-

ered sequence inconsistent with valid current evidence, such as an away-direction during a target controller hit. When the trace lacks the needed observation, the label is unknown.

Confirmation metrics use distinct event definitions. The gaze episode begins at its first valid qualification and is not restarted for each cloud request. The controller episode uses its own hold and gaze requirements. A completion event requires correct trigger selection. The evaluator records missed acknowledgment, delayed acknowledgment, or disqualification separately rather than classifying every silent interval as failure.

Advisory Model Review

The offline judge receives a bounded sequence containing current and recent state, candidate and chosen actions, actual playback, timing, and provenance. It evaluates whether the wording was useful, whether timing supported the task, whether repetition was unnecessary, and whether interruption preserved intelligibility. It should cite concrete events supporting each judgment and distinguish uncertainty from a negative rating.

The judge must not treat logs or participant text as instructions. Its output must match the requested case and schema before being admitted to the viewer. Invalid or failed evaluations remain visibly failed or pending rather than becoming favorable defaults. Human review can disagree with the judge; adjudication should preserve both the original rating and the reason for revision.

Reporting Template

For each comparison, report the recording/scenario manifest, policy revisions, number of trials, unique episodes, provider calls, latency distributions, unknown evidence, and remaining failures.

Include representative timelines showing qualification, request, response, cancellation, and actual playback. Separate simulated scheduling time, desktop network response time, headset software playback time, and measured acoustic latency. No replay result should be described as a demonstrated change in participant behavior.

APPENDIX D: STUDY READINESS AND SOURCE REGISTER

Implemented, Proposed, and Unresolved

The source snapshot documents one practice trial in the first assigned voice, two 40-trial measured blocks, nominal-angle quotas, right-trigger selection, and a final external questionnaire handoff. The protocol proposes exact pre-cue geometry and balanced cross-relative heights; the October 20 audit identifies those as implementation gaps. These statuses must be rechecked against the final release rather than inferred from the presence of future-tense method text.

Table D.1: Decisions required before confirmatory collection.

Requirement	Completion evidence
Starting geometry	Frozen origin/cross/target snapshot before cue, signed/absolute angle, time and frame provenance.
Vertical allocation	Seated-safe bounds and near-zero rule; realized 20-above/20-below schedule per block.
Power and model	Meaningful effect, theta function, random effects, simulation assumptions, participant N and enrollment ceiling.
Tracking acceptance	Independent spatial validation, unique-source timing, drift checks, and stop/recalibration rules.
Practice proficiency	Observable criterion and response if a single practice trial is insufficient.
Timeout and stopping	Prespecified comfort, safety, technical failure, and incomplete-trial treatment.
Acoustic checks	Per-voice intelligibility, comfortable levels, wording verification, and actual delivery timing.

Requirement	Completion evidence
Survey readiness	Burden/comprehension pilot, final retrospective references, validated scoring/export implementation.
Consent	Actual cloud processing, identifiable voice handling, storage, deletion, and approved participant wording.
Release freeze	Build, policy, phrase, provider and instrument versions; full-session wearer acceptance.

Source Register

The current Overleaf project is `https://www.overleaf.com/project/6a84bc097073254fbca6985b`. This first pass uses revision `f9f49e450c7b64b34c319a557114f1fa4393f003`, retrieved without overwriting the working Overleaf project. The system description preserves the current Section 3 as its foundation, with thesis-level expansion and language edits. The method and planned-analysis chapters similarly expand the current standalone sections. The 62-entry bibliography is retained as the working library; only cited entries appear in the reference list.

The full protocol handoff dated October 20 supersedes the older handoff on practice count and survey timing. The later user clarification assigns practice to the first counterbalanced voice. The repository alignment audit supplies implementation status and recorded test evidence. Historical four-condition, 14-round, dwell-selection, and immediate post-block questionnaire descriptions are not the design for this thesis.

The supplied UCF LuaLaTeX template provides the class and accessibility setup. The downloaded reference is Hoang Bao Truong's *Using Inferences from Natural Language Narratives to Improve*

Video Question Answering, UCF, 2026, available at `https://stars.library.ucf.edu/grad studies_etd_2026/199/`. Its 140-page repository PDF includes a repository cover sheet, front matter, detailed methodology, results, conclusion, appendices, and references. It is a presentation reference, not a source of findings or text for this project.

Figure Provenance

The scene illustration, procedure diagram, and guidance-memory architecture diagram were already present in the current Overleaf project and are reused here. Their editable SVG sources are included where available. They are explanatory schematics. Drawn object counts and highlight colors must not be mistaken for the frozen participant release or for an identifying target cue presented before search. The current audit describes 168 spawned objects, 21 per wall; the eight trials per nominal theta category are a different quantity.

Submission Work Remaining

This draft requires author and advisor review, a verified final study record, completed empirical analyses, and updated conclusions. Title-page administrative details and the final AI disclosure must be confirmed. Citation metadata and source-to-claim fit should receive a final audit, including any newly added quantitative literature claims. Accessibility tagging, reading order, table headers, and figure descriptions require final validation with the university's submission process. Successful compilation alone is not certification of ETD compliance or accessibility.

LIST OF REFERENCES

[1] Isayas B. Adhanom, Samantha C. Lee, Eelke Folmer, and Paul MacNeilage. GazeMetrics: An Open-Source Tool for Measuring the Data Quality of HMD-based Eye Trackers. In *ACM Symposium on Eye Tracking Research and Applications*, pages 1–5, 2020.

[2] Oscar J. Ariza N., Gerd Bruder, Nicholas Katzakis, and Frank Steinicke. Analysis of proximity-based multimodal feedback for 3d selection in immersive virtual environments. In *2018 IEEE Conference on Virtual Reality and 3D User Interfaces (VR)*, pages 327–334, 2018. doi: 10.1109/VR.2018.8446317.

[3] Oscar J. Ariza N., Markus Lange, Frank Steinicke, and Gerd Bruder. Vibrotactile assistance for user guidance towards selection targets in VR and the cognitive resources involved. In *2017 IEEE Symposium on 3D User Interfaces (3DUI)*, pages 95–98, 2017. doi: 10.1109/3DUI.2017.7893323.

[4] Dale J. Barr, Roger Levy, Christoph Scheepers, and Harry J. Tily. Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3):255–278, 2013. doi: 10.1016/j.jml.2012.11.001.

[5] Nicola Binetti, Luyan Wu, Shiping Chen, Ernst Kruijff, Simon Julier, and Duncan P. Brumby. Using visual and auditory cues to locate out-of-view objects in head-mounted augmented reality. *Displays*, 69:102032, 2021. doi: 10.1016/j.displa.2021.102032.

[6] Frank Biocca, Chad Harms, and Jenn Gregg. The Networked Minds Measure of Social Presence: Pilot Test of the Factor Structure and Concurrent Validity. In *Proceedings of the 4th Annual International Workshop on Presence*, 2001.

[7] Riccardo Bovo, Steven Abreu, Karan Ahuja, Eric J Gonzalez, Li-Te Cheng, and Mar Gonzalez-Franco. EmBARDiment: an Embodied AI Agent for Productivity in XR. In *2025 IEEE Conference Virtual Reality and 3D User Interfaces (VR)*, pages 708–717, 2025.

[8] Center for Self-Determination Theory. Intrinsic Motivation Inventory (IMI). `https://selfdeterminationtheory.org/intrinsic-motivation-inventory/`. Accessed October 20, 2026.

[9] Dongwook Choi, Taeyoon Kwon, Dongil Yang, Hyojun Kim, and Jinyoung Yeo. Designing Memory-Augmented AR Agents for Spatiotemporal Reasoning in Personalized Task Assistance. In *2025 IEEE International Symposium on Mixed and Augmented Reality Adjunct (ISMAR-Adjunct)*, pages 113–119, 2025.

[10] Alexander Diel and Michael Lewis. Deviation from typical organic voices best explains a vocal uncanny valley. *Computers in Human Behavior Reports*, 14:100430, 2024.

[11] Jan Drewes, Sascha Feder, and Wolfgang Einhäuser. Gaze During Locomotion in Virtual Reality and the Real World. *Frontiers in Neuroscience*, 15, 2021.

[12] Matt J. Dunn, Robert G. Alexander, Onyekachukwu M. Amiebenomo, Gemma Arblaster, Denize Atan, Jonathan T. Erichsen, Ulrich Ettinger, Mario E. Giardini, Iain D. Gilchrist, Ruth Hamilton, Roy S. Hessels, Scott Hodgins, Ignace T. C. Hooge, Brooke S. Jackson, Helena Lee, Stephen L. Macknik, Susana Martinez-Conde, Lee Mcilreavy, Lisa M. Muratori, Diederick C. Niehorster, Marcus Nyström, Jorge Otero-Millan, Michael M. Schlüssel, Jay E. Self, Tarkeshwar Singh, Nikolaos Smyrnis, and Andreas Sprenger. Minimal reporting guideline for research involving eye tracking (2023 edition). *Behavior Research Methods*, 56(5):4351–4357, 2024.

[13] Cathy Mengying Fang, Phoebe Chua, Samantha W. T. Chan, Joanne Leong, Andria Bao, and Pattie Maes. Leveraging AI-Generated Emotional Self-Voice to Nudge People towards their

Ideal Selves. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pages 1–20, 2025.

[14] Steve Grogorick, Michael Stengel, Elmar Eisemann, and Marcus Magnor. Subtle gaze guidance for immersive environments. In *Proceedings of the ACM Symposium on Applied Perception*, pages 1–7, 2017. doi: 10.1145/3119881.3119890.

[15] Bhasura S. Gunawardhana, Yunxiang Zhang, Qi Sun, and Zhigang Deng. Toward User-Aware Interactive Virtual Agents: Generative Multi-Modal Agent Behaviors in VR. In *2024 IEEE International Symposium on Mixed and Augmented Reality (ISMAR)*, pages 1068–1077, 2024.

[16] Siqi Guo, Minsoo Choi, Dominic Kao, and Christos Mousas. Collaborating with my Doppelgänger: The Effects of Self-similar Appearance and Voice of a Virtual Character during a Jigsaw Puzzle Co-solving Task. *Proceedings of the ACM on Computer Graphics and Interactive Techniques*, 7(1):1–23, 2024. doi: 10.1145/3651288.

[17] Sandra G. Hart and Lowell E. Staveland. Development of NASA-TLX (task load index): Results of empirical and theoretical research. In Peter A. Hancock and Najmedin Meshkati, editors, *Human Mental Workload*, volume 52 of *Advances in Psychology*, pages 139–183. Elsevier Science Publishers (North-Holland), 1988. doi: 10.1016/S0166-4115(08)62386-9.

[18] Sture Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70, 1979. JSTOR: 4615733.

[19] Dominic Kao, Rabindra Ratan, Christos Mousas, and Alejandra J. Magana. The effects of a self-similar avatar voice in educational games. *Proceedings of the ACM on Human-Computer Interaction*, 5(CHI PLAY):1–28, 2021. doi: 10.1145/3474665.

[20] Jieun Kim and Hayeon Song. My Voice as a Daily Reminder: Self-Voice Alarm for Daily Goal Achievement. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pages 1–16, 2024.

[21] Tiffany C. K. Kwok, Peter Kiefer, and Martin Raubal. Two-step gaze guidance. In *Proceedings of the 2022 International Conference on Multimodal Interaction*, pages 299–309, 2022. doi: 10.1145/3536221.3556612.

[22] Daniel Lange, Tim Claudius Stratmann, Uwe Gruenefeld, and Susanne Boll. HiveFive: Immersion preserving attention guidance in virtual reality. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, pages 1–13, 2020. doi: 10.1145/3313831.3376803.

[23] Kwan-Min Lee and Clifford Nass. Social-Psychological Origins of Feelings of Presence: Creating Social Presence With Machine-Generated Voices. *Media Psychology*, 7(1):31–45, 2005.

[24] Chenyi Li, Guande Wu, Gromit Yeuk-Yin Chan, Dishita Gdi Turakhia, Sonia Castelo Quispe, Dong Li, Leslie Welch, Claudio Silva, and Jing Qian. Satori: Towards Proactive AR Assistant with Belief-Desire-Intention User Modeling. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, pages 1–24, 2025.

[25] Viktor Losing, Lukas Rottkamp, Michael Zeunert, and Thies Pfeiffer. Guiding visual search tasks using gaze-contingent auditory feedback. In *Proceedings of the 2014 ACM International Joint Conference on Pervasive and Ubiquitous Computing: Adjunct Publication*, pages 1093–1102, 2014. doi: 10.1145/2638728.2641687.

[26] Edward McAuley, Terry Duncan, and Vance V. Tammen. Psychometric Properties of the Intrinsic Motivation Inventory in a Competitive Sport Setting: A Confirmatory Factor Analysis. *Research Quarterly for Exercise and Sport*, 60(1):48–58, 1989.

[27] Sara Moussawi and Marios Koufaris. Perceived Intelligence and Perceived Anthropomorphism of Personal Intelligent Agents: Scale Development and Validation. In *Proceedings of the Annual Hawaii International Conference on System Sciences*, 2019.

[28] Clifford Nass and Kwan Min Lee. Does computer-synthesized speech manifest personality? Experimental tests of recognition, similarity-attraction, and consistency-attraction. *Journal of Experimental Psychology: Applied*, 7(3):171–181, 2001.

[29] Clifford Nass and Jonathan Steuer. Voices, Boxes, and Sources of Messages: Computers and Social Actors. *Human Communication Research*, 19(4):504–527, 1993.

[30] Diederick C. Niehorster, Marcus Nyström, Roy S. Hessels, Jeroen S. Benjamins, Richard Andersson, and Ignace T. C. Hooge. The fundamentals of eye tracking, Part 7: Determining data quality. *Behavior Research Methods*, 58(7), 2026.

[31] Richard M. Ryan. Control and information in the intrapersonal sphere: An extension of cognitive evaluation theory. *Journal of Personality and Social Psychology*, 43(3):450–461, 1982.

[32] Dario D. Salvucci and Joseph H. Goldberg. Identifying fixations and saccades in eye-tracking protocols. In *Proceedings of the symposium on Eye tracking research & applications - ETRA '00*, pages 71–78, 2000.

[33] Immo Schuetz and Katja Fiehler. Eye tracking in virtual reality: Vive pro eye spatial accuracy, precision, and calibration reliability. *Journal of Eye Movement Research*, 15(3), 2022.

[34] Alexandra Sipatchin, Siegfried Wahl, and Katharina Rifai. Eye-Tracking for Clinical Ophthalmology with Virtual Reality (VR): A Case Study of the HTC Vive Pro Eye's Usability. *Healthcare*, 9(2):180, 2021.

[35] Eugene C. Snyder, Sanjana Mendu, S. Shyam Sundar, and Saeed Abdullah. Busting the one-voice-fits-all myth: Effects of similarity and customization of voice-assistant personality. *International Journal of Human-Computer Studies*, 180:103126, 2023.

[36] Niklas Stein, Tamara Watson, Markus Lappe, Maren Westendorf, and Szonya Durant. Eye and head movements in visual search in the extended field of view. *Scientific Reports*, 14, 2024. doi: 10.1038/s41598-024-59657-5.

[37] Amelia C. Warden, Christopher D. Wickens, Domenick Mifsud, Shannon Ourada, Benjamin A. Clegg, and Francisco R. Ortega. Visual search in augmented reality: Effect of target cue type and location. *Proceedings of the Human Factors and Ergonomics Society Annual Meeting*, 66(1):373–377, 2022. doi: 10.1177/1071181322661260.

[38] Yichen Yu, Yifan Jiang, Mandy Lui, and Qiao Jin. GenLARP: Enabling Immersive Live Action Role-Play through LLM-Generated Worlds and Characters. In *2025 IEEE International Symposium on Mixed and Augmented Reality Adjunct (ISMAR-Adjunct)*, pages 763–764, 2025.

[39] Johannes Zagermann, Ulrike Pfeil, and Harald Reiterer. Measuring Cognitive Load using Eye Tracking Technology in Visual Computing. In *Proceedings of the Sixth Workshop on Beyond Time and Errors on Novel Evaluation Methods for Visualization*, pages 78–85, 2016.

[40] Yuchong Zhang, Adam Nowak, Yueming Xuan, Andrzej Romanowski, and Morten Fjeld. See or hear? exploring the effect of visual/audio hints and gaze-assisted instant post-task feedback for visual search tasks in AR. In *2023 IEEE International Symposium on Mixed and Augmented Reality (ISMAR)*, pages 1113–1122, 2023. doi: 10.1109/ISMAR59233.2023.00128.