

Major Papers Behind Gaze Contingency

What the core visual-search, gaze, proximity-feedback, workload, and self-similar voice papers studied - and exactly what each one can defend in our experiment.

23 major papers

Selected for direct methodological value, not merely topical similarity.

Five evidence families

Search task, cue design, eye/head measurement, proximity feedback, and voice identity.

Decision-oriented summaries

Every entry states what transfers and what does not.

Executive Read

The literature supports a controlled seated object-search study, but it does not validate every detail of the current implementation. The strongest defensible design is a balanced 2 x 2 experiment that separates gaze contingency from voice identity, uses controlled conjunction-search layouts, matches hint opportunities, and records both trial outcomes and frame-level eye/head behavior.

Evidence	What the literature supports	Implication
Search field	Seated broad-field and 360-degree search are established paradigms.	Start with a controlled 180-degree field; treat 360 degrees as a later extension.
Hints	Spatial, visual, audio, and gaze-linked guidance can improve search.	Match timing and exposure so gaze awareness is the only systematic difference.
Proximity	Proximity aids retrieval and target acquisition, but continuous proximity speech is not well validated for visual search.	Use discrete, coarse-to-fine region states with cooldown and hysteresis.
Measurement	Eye and head behavior provide complementary evidence; peripheral processing exceeds a single ray.	Keep raw gaze/head streams and label XRI hover as an inspection proxy until validated.
Voice	Self-similar voice can increase identification, likability, and trust, but also eeriness.	Hold wording/prosody constant and measure similarity, authenticity, and discomfort.

Bottom line: the papers justify the research question and core measures. They do not justify copying another headset's fixation thresholds, treating hover as fixation, or comparing useful gaze-aware hints against empty encouragement.

Closest direct precedents

These papers most closely resemble the target task or manipulate the same guidance problem. They should carry the greatest weight in the methods defense.

DIRECT

[Visual Search in Augmented Reality: Effect of Target Cue Type and Location](#)

Warden et al. | HFES 2022 | Cited by 19 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A seated HoloLens 2 search task spanning roughly 180 degrees, with 32 physical objects and four cue conditions: a world-referenced arrow, target icon, minimap, and no cue. Gaze indicated the inspected object before explicit confirmation.

Main result. Spatially conformal cues, especially the arrow, improved far-domain search. Central cue placement reduced access effort, although visually dense overlays could create clutter.

What we can borrow. The strongest precedent for a seated semicircular object field, explicit response confirmation, target-cue comparisons, completion time, and selection errors.

Transfer caution. Its 'display proximity' means cue-to-referent closeness, not warmer-colder distance feedback.

DIRECT

[See or Hear? Visual/Audio Hints and Gaze-Assisted Feedback for Visual Search in AR](#)

Zhang et al. | IEEE ISMAR 2023 | Cited by 9 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Two HMD-based AR studies with 96 total participants searching for books while receiving visual hints, audio hints, combined hints, and gaze-assisted post-task feedback.

Main result. Visual and audio hints could improve search efficiency, but their combination was conditionally beneficial. Gaze-assisted feedback produced clearer benefits than simply adding more cue channels.

What we can borrow. Primary support for spoken hints, controlled target-search timing, familiarization, NASA-TLX, and gaze-linked feedback in AR.

Transfer caution. The task and feedback were book-search specific; exact scripts and thresholds require adaptation.

DIRECT

Searching Across Realities: ERPs and Eye-Tracking Correlates of Visual Search in MR

Chiossi et al. | IEEE ISMAR / TVCG | Cited by 12 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A within-participant visual-search comparison across augmented reality and augmented virtuality, with physical and virtual targets. The study combines behavioral outcomes, EEG, fixation/saccade features, pupil-based workload, and NASA-TLX.

Main result. Search was more efficient in augmented virtuality; AR produced more scattered gaze and greater workload. Virtual targets were generally processed more efficiently than physical targets.

What we can borrow. The central source for separating performance, gaze allocation, target type, and workload; also supports calibration and criterion-based practice.

Transfer caution. Its eye-event thresholds were created for different hardware and sampling behavior; do not copy them directly to Focus Vision logs.

DIRECT

Investigating Search Among Physical and Virtual Objects Under Different Lighting Conditions

Kim et al. | IEEE ISMAR / TVCG | Cited by 34 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Forty-eight HoloLens 2 users searched an outdoor courtyard for virtual gems among physical and virtual scene objects. Lighting and a secondary auditory-monitoring task manipulated environmental and cognitive demand.

Main result. Lighting, physical-versus-virtual object status, and dual-task load changed visibility, discovery, recall, gaze, and movement behavior.

What we can borrow. Supports gaze/head/path logging, target occlusion, cognitive-load manipulation, object recall, and explicit treatment of physical and virtual content.

Transfer caution. Its walking, outdoor treasure hunt has much more locomotion and environmental variability than a seated laboratory task.

DIRECT

The Impact of Navigation Aids on Search Performance and Object Recall in Wide-Area AR

Kumaran et al. | ACM CHI 2023 | Cited by 35 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A 24-participant outdoor AR search study comparing no aid with a compass, radar, and in-world vertical arrows while measuring search, movement, gaze, surveys, and incidental object recall.

Main result. All navigation aids improved search, with the strongest preference for in-world arrows. Assistance also shifted attention and could reduce awareness of surrounding physical objects.

What we can borrow. Use search efficiency and environmental awareness together so guidance is not evaluated only by speed.

Transfer caution. Navigation aids reveal location more directly than our proposed conversational guidance.

DIRECT

Cuing Multiple Targets for Visual Search in Virtual Reality

Kelley et al. | IEEE VRW 2025 | Cited by 0 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A 360-degree VR visual-search task with multiple targets and distractors, comparing gaze-line, arrow/wedge, and no-cue conditions.

Main result. Explicit spatial cueing markedly improved target localization; gaze-line guidance was especially accurate and fast in the reported study.

What we can borrow. Supports balancing initially out-of-view targets, measuring angular search behavior, and retaining a no-cue or noncontingent baseline.

Transfer caution. This is a short workshop publication and evaluates visual cues, not a spoken gaze-aware agent.

Task and cue structure

These studies explain how field of view, distractor density, cue form, and cue reliability change search behavior.

GUIDANCE

Subtle Cueing for Visual Search in Augmented Reality

Lu, Duh, and Feiner | IEEE ISMAR 2012 | Cited by 43 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Two studies of low-opacity visual modulation intended to attract attention toward targets in cluttered outdoor imagery without adding an obvious arrow or overlay.

Main result. Subtle guidance can improve search within an appropriate range of scene clutter and cue visibility, but its effectiveness depends on cue parameters.

What we can borrow. Supports assistance that narrows search without explicitly revealing the answer.

Transfer caution. Subtle visual modulation and spoken guidance operate through different channels.

DIRECT

The Effect of Narrow FOV and Information Density on Visual Search Performance in AR

Trepkowski et al. | IEEE VR 2019 | Cited by 60 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A series of visual-search experiments that manipulated effective display field of view and information density using a HoloLens and a head-tracked tiled display.

Main result. Larger effective FOV improved search by roughly 7-28 percent depending on condition. Density and limited peripheral access materially constrained search.

What we can borrow. Control set size, object spacing, eccentricity, and rendered angular size; treat headset FOV as part of the task rather than background hardware.

Transfer caution. The stimuli were labels and symbols, not volumetric objects in passthrough MR.

DIRECT

Measuring Visual Search and Distraction in Immersive Virtual Reality

Olk et al. | Royal Society Open Science | Cited by 72 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Participants searched a virtual kitchen countertop for everyday targets among distractors while target discriminability and peripheral-flanker congruence were manipulated.

Main result. Immersive VR reproduced expected visual-search and distraction effects, demonstrating that naturalistic 3D object arrangements can remain experimentally controlled.

What we can borrow. Supports conjunction-like object search, controlled distractor similarity, reaction time, and peripheral distraction.

Transfer caution. The scene was frontal and compact rather than a broad egocentric search field.

GUIDANCE

Comparing Non-Visual and Visual Guidance for Narrow-FOV AR Displays

Marquardt et al. | IEEE ISMAR / TVCG | Cited by 56 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Three studies compare audio-tactile guidance with the EyeSee360 visual technique for locating out-of-view objects under a simulated narrow-FOV AR setup.

Main result. Visual guidance was generally faster, while audio-tactile guidance preserved hit rate and improved situation awareness by reducing visual-channel competition.

What we can borrow. Supports nonvisual guidance when the search task already loads the visual channel; include situation awareness and workload, not only speed.

Transfer caution. The AR environment was simulated in VR, and audio-tactile direction cues are not equivalent to semantic speech.

GUIDANCE

Exploring the Effect of Visual Cues on Eye Gaze During AR-Guided Tasks

Seeliger et al. | IEEE ISMAR Adjunct | Cited by 26 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. An AR picking-and-assembly study comparing cue presentation patterns, including in-view/out-of-view, static/dynamic, and sequential/simultaneous guidance, while recording gaze behavior.

Main result. Cue design changed gaze distribution, duration, and scanpath behavior, demonstrating that guidance can restructure attention even when task success improves.

What we can borrow. Log gaze around each hint and compare scanpath changes, not just final completion.

Transfer caution. Procedural picking is different from open visual search.

GUIDANCE

Fast, Accurate, but Sometimes Too-Compelling Support

Warden et al. | IEEE THMS | Cited by 17 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Fifty-six participants searched a 3D room containing 48 objects while receiving arrow, minimap, icon, or no guidance. For one group the automated cue was wrong on 17 percent of trials.

Main result. Location cues improved speed and accuracy when correct but created larger costs and automation bias when wrong.

What we can borrow. Log hint validity, compliance, recovery, and erroneous selections. Treat inaccurate guidance as a measurable system event.

Transfer caution. Do not intentionally introduce unreliable hints unless reliability is a formal, powered factor.

Gaze, head movement, and workload

These papers justify recording eye and head behavior separately and treating workload as more than completion time.

MEASUREMENT

The Importance of Peripheral Vision When Searching 3D Real-World Scenes

David et al. | Journal of Vision | Cited by 39 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A gaze-contingent VR study that restricted central or peripheral scene information with windows and scotomas during search through 3D scenes.

Main result. Peripheral information plays a major role in guiding search and cannot be represented by a single infinitesimal gaze ray.

What we can borrow. Use an angular tolerance or object-neighborhood interpretation for attention; distinguish direct hover from likely peripheral processing.

Transfer caution. The paper manipulates visual availability rather than validating a specific interaction cone for our headset.

MEASUREMENT

Eye and Head Movements in Visual Search in the Extended Field of View

Stein et al. | Scientific Reports | Cited by 13 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A seated extended-field VR search study measuring how eye movements and head turns jointly bring targets into view and support inspection.

Main result. Eye and head behavior contribute differently across search stages, making combined gaze direction alone insufficient for explaining search strategy.

What we can borrow. Analyze eye latency, head latency, total yaw, sector transitions, and eye-versus-head contribution separately.

Transfer caution. Derive formulas from the signals actually exported by Focus Vision rather than assuming equivalent native samples.

MEASUREMENT

Recognizing User Tasks From Eye and Head Movements in Immersive VR

EHTask authors | IEEE TVCG | Cited by 58 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Eye and head movement features were used to distinguish immersive-VR activities, including visual search.

Main result. Eye and head streams contain complementary information about task state and user behavior.

What we can borrow. Retain both frame-level gaze and head direction for later analysis and replay.

Transfer caution. Task classification is not the same as measuring workload or search quality.

Proximity and adaptive hints

These studies define several different meanings of proximity. Only NeighboAR is an object-retrieval task; Ariza and Hu concern final target acquisition.

PROXIMITY

NeighboAR: Proximity- and Gaze-Based Object Grouping for AR Retrieval

Slavuljica et al. | PACMHCI ETRA 2024 | Cited by 1 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Seventeen participants inspected 30 grocery products across four cabinets. The system built a graph connecting objects within 30 cm and later displayed the target plus the nearby object with the longest prior gaze dwell.

Main result. The contextual landmark reduced retrieval time by 33 percent, errors by 71 percent, and perceived workload by about 10 percent.

What we can borrow. Represent the object field as spatial neighborhoods and provide group-level rather than exact-target hints.

Transfer caution. This is memory-supported retrieval after inspection, not continuous warmer-colder search.

PROXIMITY

Analysis of Proximity-Based Multimodal Feedback for 3D Selection

Ariza et al. | IEEE VR 2018 | Cited by 38 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A Fitts' Law experiment comparing binary and continuous visual, auditory, tactile, and bimodal feedback whose intensity depended on cursor-to-target distance.

Main result. Feedback type affected both the ballistic approach and final correction phases of 3D selection and significantly changed performance.

What we can borrow. Separate approach guidance from confirmation and define the distance-to-feedback function explicitly.

Transfer caution. This is final target acquisition, not finding an unknown target among distractors.

PROXIMITY

Optimizing Moving Target Selection in VR With Proximity-Based Feedback

Hu et al. | IEEE VR 2025 | Cited by 5 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Two studies of moving-target selection cross visual, audio, and haptic modalities with binary, continuous, and partial proximity mappings. Target width and speed were also varied.

Main result. Partial visual feedback combined fast approach information with clear contact confirmation. Binary audio and haptic feedback were generally preferred to continuous signals, which could be uncomfortable or hard to interpret near the target.

What we can borrow. Use coarse-to-fine visual/verbal states plus a distinct confirmation event; avoid continuous speech or vibration on every gaze movement.

Transfer caution. The moving target was always in front of the participant, specifically excluding out-of-view search.

Self-similar voice and intelligent agents

This evidence supports the voice factor and manipulation checks, while showing why cloned voice quality and eeriness cannot be ignored.

VOICE

Collaborating With My Doppelganger

Guo et al. | PACM Computer Graphics and Interactive Techniques | Cited by 20 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Twenty-five participants solved collaborative VR jigsaw puzzles in a 2 x 2 design crossing self-similar appearance and cloned self-similar voice.

Main result. Voice similarity increased agent likability and believability but could also contribute to eeriness.

What we can borrow. Primary precedent for the factorial voice manipulation, standardized cloning workflow, counterbalancing, and post-condition similarity checks.

Transfer caution. The voice belonged to a collaborative avatar, not a gaze-monitoring guide agent.

VOICE

The Effects of a Self-Similar Avatar Voice in Educational Games

Kao et al. | ACM CHI PLAY 2021 | Cited by 36 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. A 2 x 3 study crossed similar versus dissimilar avatar voice with pitch manipulation during a Java-programming game.

Main result. Self-similar voice increased performance, time spent, identification, competence, relatedness, and immersion; identification mediated several outcomes.

What we can borrow. Supports hypotheses about identification and engagement and motivates measuring competence and relatedness alongside performance.

Transfer caution. An educational game is motivationally different from repeated laboratory search trials.

VOICE

AI-Determined Similarity Increases Likability and Trustworthiness of Human Voices

Jaggy et al. | PLOS ONE | Cited by 8 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Five preregistered online experiments validated embedding-based voice similarity; one experiment compared 100 same-gender voices against each participant's recorded voice.

Main result. Acoustically self-similar voices received modestly higher likability and trustworthiness ratings.

What we can borrow. Provides a measurable voice-similarity mechanism and supports a continuous manipulation check rather than assuming cloning worked.

Transfer caution. It demonstrates social evaluation, not improved XR task performance.

VOICE

Perception and Social Evaluation of Cloned and Recorded Voices

McGettigan et al. | CHB: Artificial Humans | Cited by 10 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Familiar pairs and unfamiliar listeners judged cloned and recorded self/friend voices on authenticity and social traits including trustworthiness, attractiveness, competence, and dominance.

Main result. Familiar listeners were sensitive to clone authenticity, and perceived cloning accuracy predicted social evaluation. A friend's voice could be preferred to one's own.

What we can borrow. Measure similarity, authenticity, naturalness, and preference separately; record synthesis failures and provider settings.

Transfer caution. A technically poor clone can reverse the intended manipulation.

AGENT

Comparing Referential Interactions With an Intelligent Assistant in VR

Kaschub et al. | IEEE VR 2026 | Cited by 0 (OpenAlex snapshot, Aug. 5, 2026)

What it describes. Thirty-nine users referred to virtual objects while interacting with an intelligent assistant through verbal description, gaze, or pointing.

Main result. The work treats gaze as a low-effort referential signal and evaluates task time, prompting, NASA-TLX, and user experience.

What we can borrow. Supports gaze as implicit context for an assistant and motivates explicit comparison of effort, prompts, and perceived control.

Transfer caution. Referencing a known object is not the same as receiving target-search guidance.

Experiment Defense Matrix

Use this table when explaining why each major procedural choice exists. The papers support the design principle; our exact thresholds, layouts, and timing still require piloting on the Vive Focus Vision.

Project choice	Primary citations	Defense
Seated 180-degree field	Warden 2022	Direct precedent: seated HoloLens search over 32 objects across a broad frontal field.
Conjunction distractors	Olk 2018; Trepkowski 2019	Controlled target-distractor similarity and density produce measurable search difficulty.
Spoken hints	Zhang 2023; Marquardt 2020	Audio can guide AR search without adding visual clutter, but must be evaluated for workload.
Gaze-aware policy	Zhang 2023; Kaschub 2026	Gaze can provide low-effort context for feedback and referential interaction.
Coarse spatial neighborhoods	NeighboAR 2024	Nearby, well-inspected objects can act as useful contextual landmarks.
Coarse-to-fine feedback	Ariza 2018; Hu 2025	Approach information and final confirmation should be distinct feedback stages.
Peripheral tolerance	David 2021	A single ray underrepresents information available in peripheral vision.
Separate eye/head metrics	Stein 2024; EHTask 2021	Eye and head streams make different contributions to extended-field search.
Workload measurement	Chiossi 2024; Zhang 2023	NASA-TLX and gaze/pupil measures complement performance outcomes.
Incorrect-hint logging	Warden 2023	Powerful cues can create automation bias and larger costs when wrong.
Self-similar voice factor	Guo 2024; Kao 2021	Voice similarity can alter identification, engagement, likability, and task behavior.
Voice manipulation checks	Jaggy 2025; McGettigan 2025	Perceived similarity and authenticity must be measured, not assumed.

Recommended Evidence Stack

- Lead with **Warden 2022** for the seated broad-field layout and cue comparison.
- Use **Zhang 2023** for spoken AR search hints and gaze-linked feedback.
- Use **Chiossi 2024** plus **David 2021** and **Stein 2024** for gaze, peripheral processing, head motion, and workload.
- Use **NeighboAR 2024**, **Ariza 2018**, and **Hu 2025** to motivate coarse spatial neighborhoods and staged proximity feedback.
- Use **Guo 2024** as the main self-similar voice precedent, supported by **Kao 2021**, **Jaggy 2025**, and **McGettigan 2025**.
- Use **Warden 2023** to justify logging hint errors, compliance, and recovery rather than assuming the assistant is always correct.

References and Source Notes

Links resolve to DOI landing pages. Citation counts are comparative OpenAlex values recorded in the repository inventory on August 5, 2026; they are not definitive impact measures and will change over time.

1. Warden et al. (2022). [Visual Search in Augmented Reality: Effect of Target Cue Type and Location](https://doi.org/10.1177/1071181322661260). HFES 2022. <https://doi.org/10.1177/1071181322661260>
2. Zhang et al. (2023). [See or Hear? Visual/Audio Hints and Gaze-Assisted Feedback for Visual Search in AR](https://doi.org/10.1109/ISMAR59233.2023.00128). IEEE ISMAR 2023. <https://doi.org/10.1109/ISMAR59233.2023.00128>
3. Chiossi et al. (2024). [Searching Across Realities: ERPs and Eye-Tracking Correlates of Visual Search in MR](https://doi.org/10.1109/TVCG.2024.3456172). IEEE ISMAR / TVCG. <https://doi.org/10.1109/TVCG.2024.3456172>
4. Kim et al. (2022). [Investigating Search Among Physical and Virtual Objects Under Different Lighting Conditions](https://doi.org/10.1109/TVCG.2022.3203093). IEEE ISMAR / TVCG. <https://doi.org/10.1109/TVCG.2022.3203093>
5. Kumaran et al. (2023). [The Impact of Navigation Aids on Search Performance and Object Recall in Wide-Area AR](https://doi.org/10.1145/3544548.3581413). ACM CHI 2023. <https://doi.org/10.1145/3544548.3581413>
6. Kelley et al. (2025). [Cuing Multiple Targets for Visual Search in Virtual Reality](https://doi.org/10.1109/VRW66409.2025.00257). IEEE VRW 2025. <https://doi.org/10.1109/VRW66409.2025.00257>
7. Lu, Duh, and Feiner (2012). [Subtle Cueing for Visual Search in Augmented Reality](https://doi.org/10.1109/ISMAR.2012.6402553). IEEE ISMAR 2012. <https://doi.org/10.1109/ISMAR.2012.6402553>
8. Trepkowski et al. (2019). [The Effect of Narrow FOV and Information Density on Visual Search Performance in AR](https://doi.org/10.1109/VR.2019.8798312). IEEE VR 2019. <https://doi.org/10.1109/VR.2019.8798312>
9. Olk et al. (2018). [Measuring Visual Search and Distraction in Immersive Virtual Reality](https://doi.org/10.1098/rsos.172331). Royal Society Open Science. <https://doi.org/10.1098/rsos.172331>
10. Marquardt et al. (2020). [Comparing Non-Visual and Visual Guidance for Narrow-FOV AR Displays](https://doi.org/10.1109/TVCG.2020.3023605). IEEE ISMAR / TVCG. <https://doi.org/10.1109/TVCG.2020.3023605>
11. Seeliger et al. (2021). [Exploring the Effect of Visual Cues on Eye Gaze During AR-Guided Tasks](https://doi.org/10.1109/ISMAR-Adjunct54149.2021.00041). IEEE ISMAR Adjunct. <https://doi.org/10.1109/ISMAR-Adjunct54149.2021.00041>
12. Warden et al. (2023). [Fast, Accurate, but Sometimes Too-Compelling Support](https://doi.org/10.1109/THMS.2023.3302152). IEEE THMS. <https://doi.org/10.1109/THMS.2023.3302152>
13. David et al. (2021). [The Importance of Peripheral Vision When Searching 3D Real-World Scenes](https://doi.org/10.1167/jov.21.7.3). Journal of Vision. <https://doi.org/10.1167/jov.21.7.3>
14. Stein et al. (2024). [Eye and Head Movements in Visual Search in the Extended Field of View](https://doi.org/10.1038/s41598-024-59657-5). Scientific Reports. <https://doi.org/10.1038/s41598-024-59657-5>
15. EHTask authors (2021). [Recognizing User Tasks From Eye and Head Movements in Immersive VR](https://doi.org/10.1109/TVCG.2021.3138902). IEEE TVCG. <https://doi.org/10.1109/TVCG.2021.3138902>
16. Slavuljica et al. (2024). [NeighboAR: Proximity- and Gaze-Based Object Grouping for AR Retrieval](https://doi.org/10.1145/3655599). PACMHCI ETRA 2024. <https://doi.org/10.1145/3655599>
17. Ariza et al. (2018). [Analysis of Proximity-Based Multimodal Feedback for 3D Selection](https://doi.org/10.1109/VR.2018.8446317). IEEE VR 2018. <https://doi.org/10.1109/VR.2018.8446317>
18. Hu et al. (2025). [Optimizing Moving Target Selection in VR With Proximity-Based Feedback](https://doi.org/10.1109/VR59515.2025.00030). IEEE VR 2025. <https://doi.org/10.1109/VR59515.2025.00030>
19. Guo et al. (2024). [Collaborating With My Doppelganger](https://doi.org/10.1145/3651288). PACM Computer Graphics and Interactive Techniques. <https://doi.org/10.1145/3651288>
20. Kao et al. (2021). [The Effects of a Self-Similar Avatar Voice in Educational Games](https://doi.org/10.1145/3474665). ACM CHI PLAY 2021. <https://doi.org/10.1145/3474665>
21. Jaggy et al. (2025). [AI-Determined Similarity Increases Likability and Trustworthiness of Human Voices](https://doi.org/10.1371/journal.pone.0318890). PLOS ONE. <https://doi.org/10.1371/journal.pone.0318890>
22. McGettigan et al. (2025). [Perception and Social Evaluation of Cloned and Recorded Voices](https://doi.org/10.1016/j.chbah.2025.100143). CHB: Artificial Humans. <https://doi.org/10.1016/j.chbah.2025.100143>
23. Kaschub et al. (2026). [Comparing Referential Interactions With an Intelligent Assistant in VR](https://doi.org/10.1109/VR67842.2026.00066). IEEE VR 2026. <https://doi.org/10.1109/VR67842.2026.00066>

Scope note: This is a decision-oriented major-paper brief, not a systematic review. The broader repository inventory contains additional direct, transferable, adjacent, clinical, and information-retrieval work.